



# INNOVATE2018

## ONLINE CONFERENCE



# Supercharge Your Apps with Amazon Neptune Graph Database (Level 200)

Clayton Brown, Solution Architect

# Fully Managed Database Options on AWS

## Relational Databases



Amazon RDS



Amazon Redshift

Aurora

Commercial

Community

Data Warehouse



ORACLE



PostgreSQL



PostgreSQL



AWS Database Migration Service

## Non-Relational Databases



Amazon  
DynamoDB



Amazon  
ElastiCache



Amazon  
Neptune

Key Value

In-Memory  
Data Store

Graph


Document



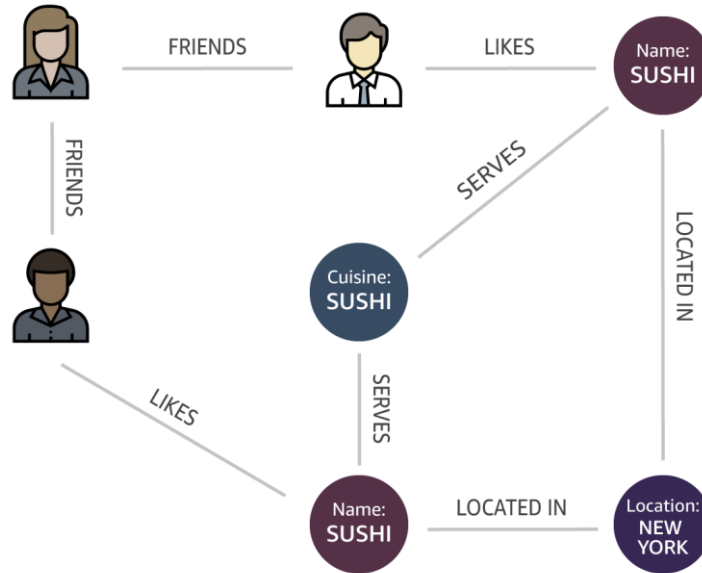
# Agenda

- Building Applications on Highly Connected Data
- Types of Graphs and How to Query Them
- Property Graph and Apache TinkerPop Friend Recommendation Example
- RDF Knowledge Graph Example
- Overview of Amazon Neptune's Fully Managed, Enterprise-Ready Features

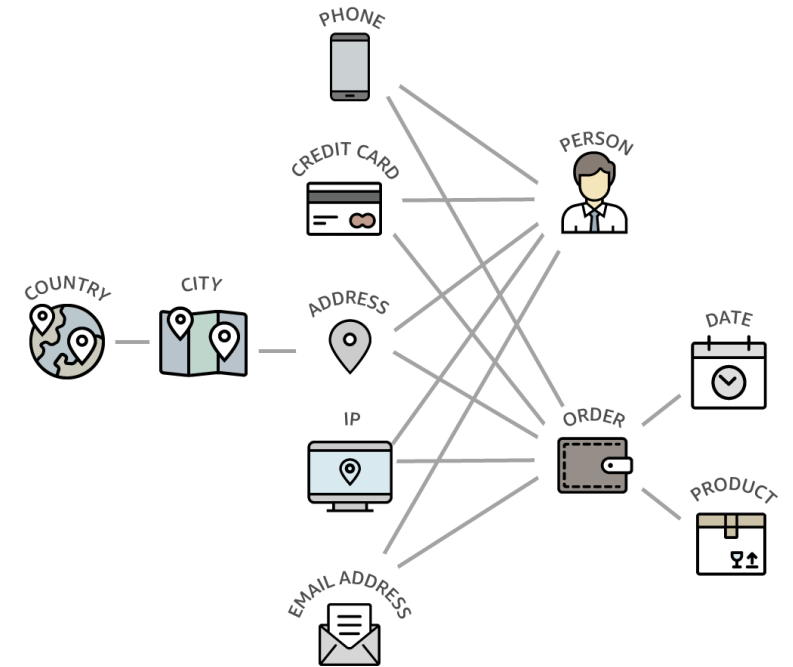
# Highly Connected Data



Social Networks



Restaurant Recommendations




Retail Fraud Detection

# Use Cases for Highly Connected Data




Social Networking




Recommendations




Knowledge Graphs



Fraud Detection

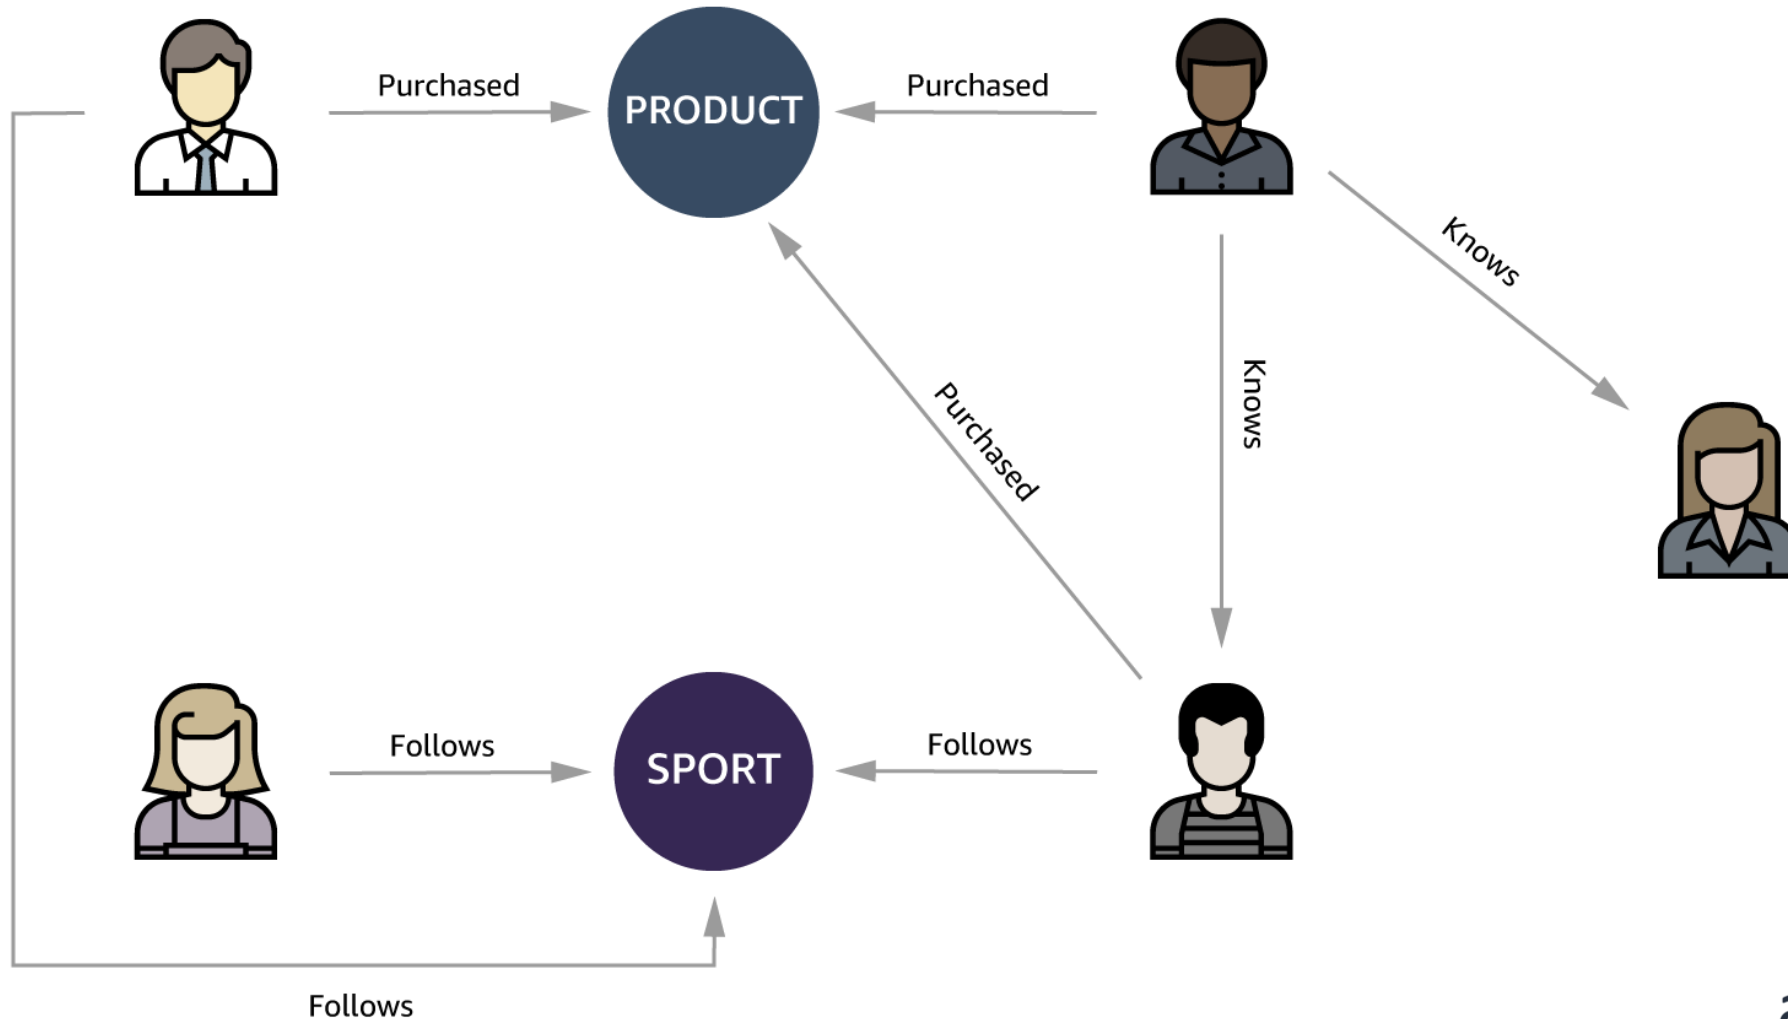


Life Sciences

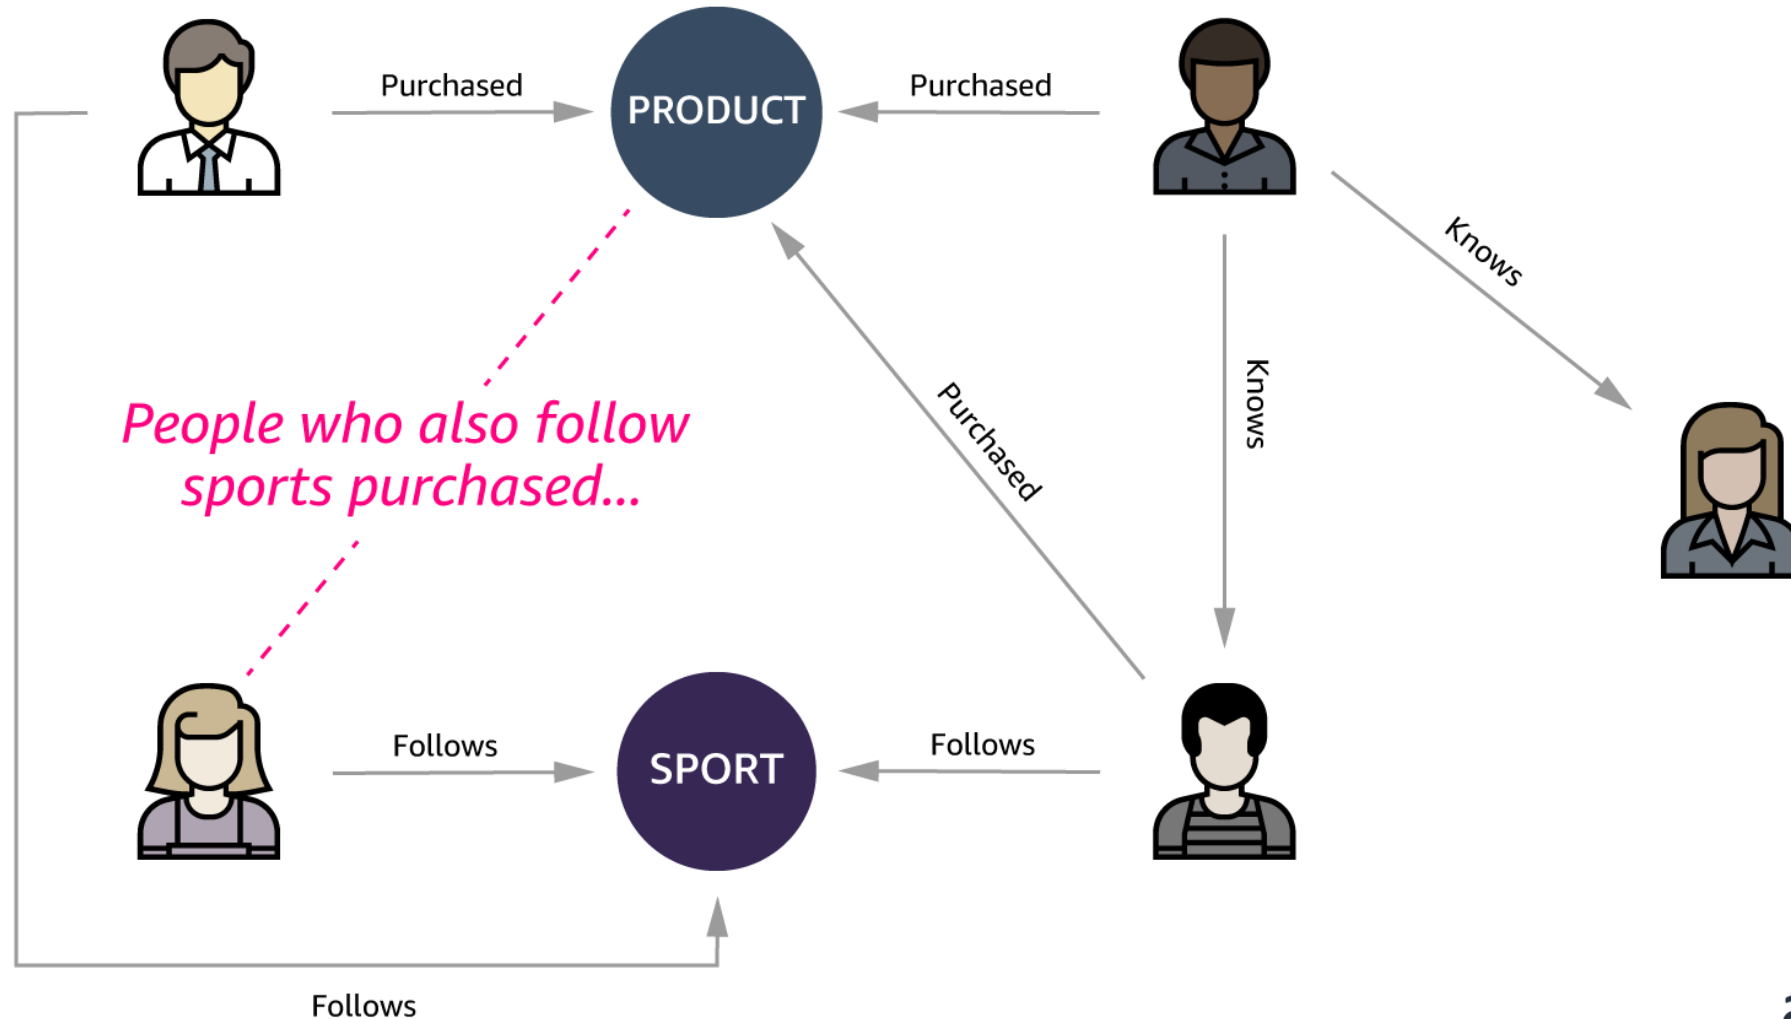


Network & IT Operations

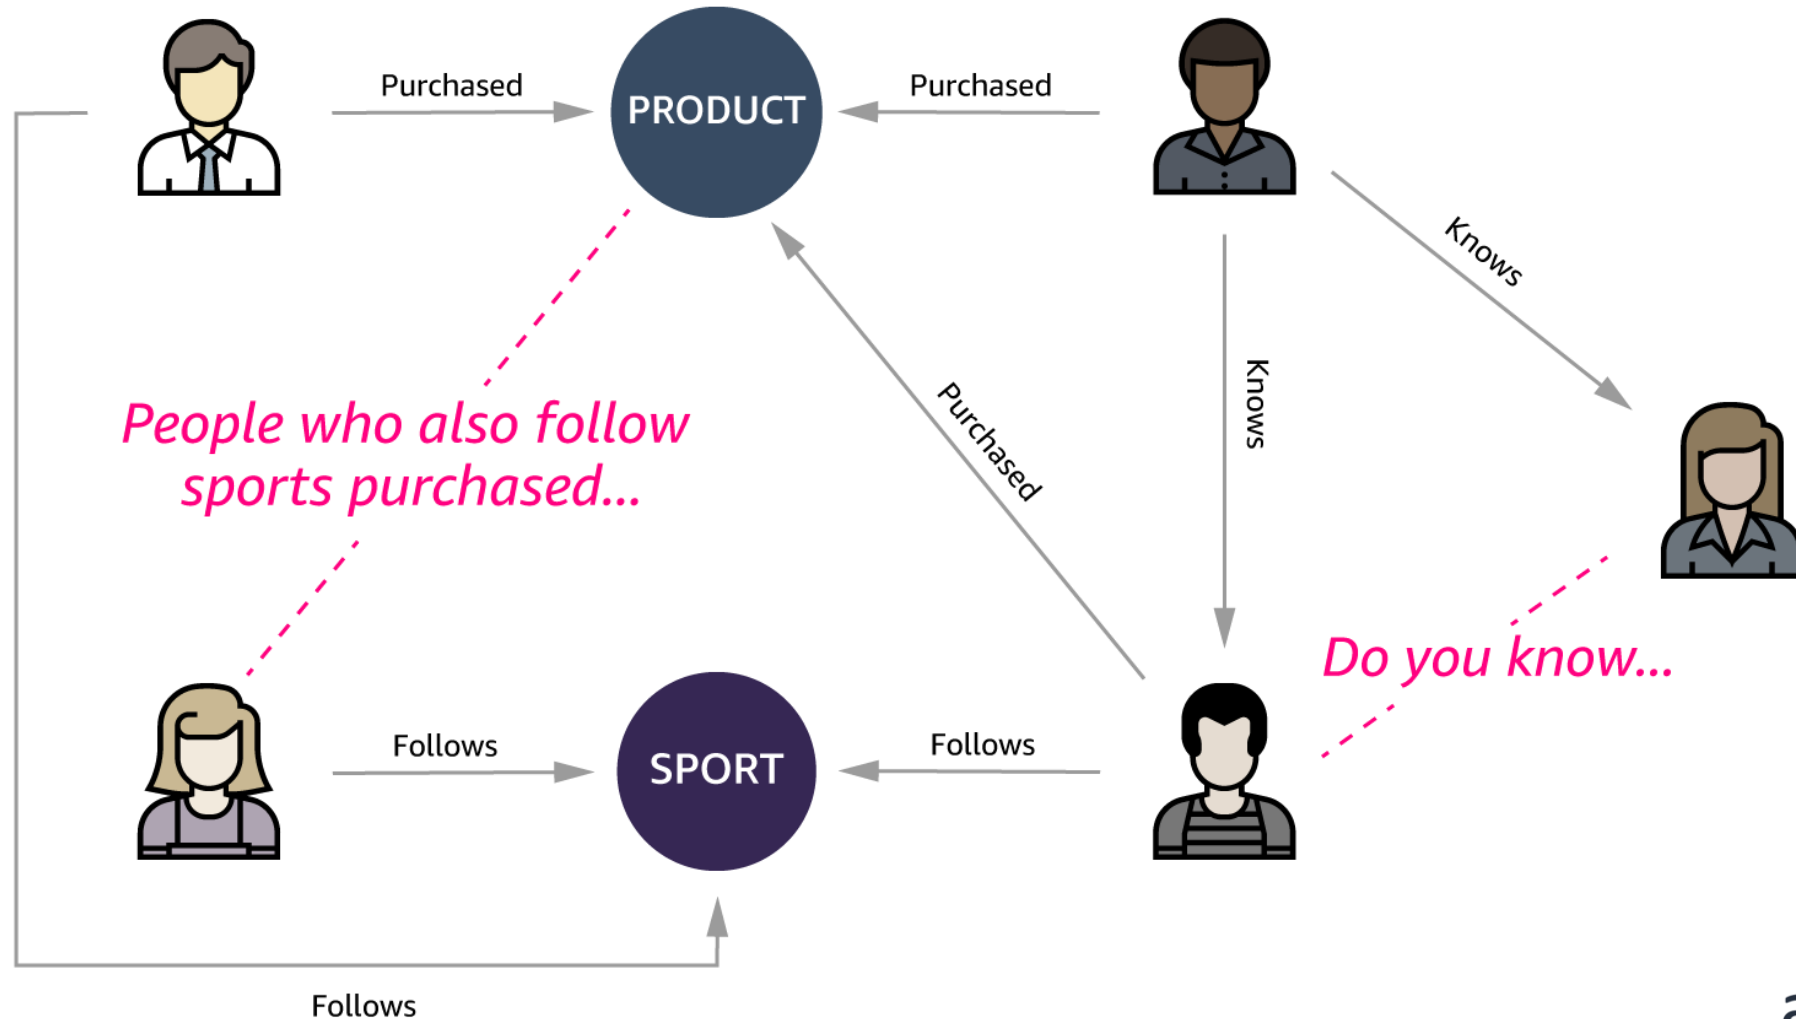
# Recommendations Based on Relationships



# Recommendations Based on Relationships

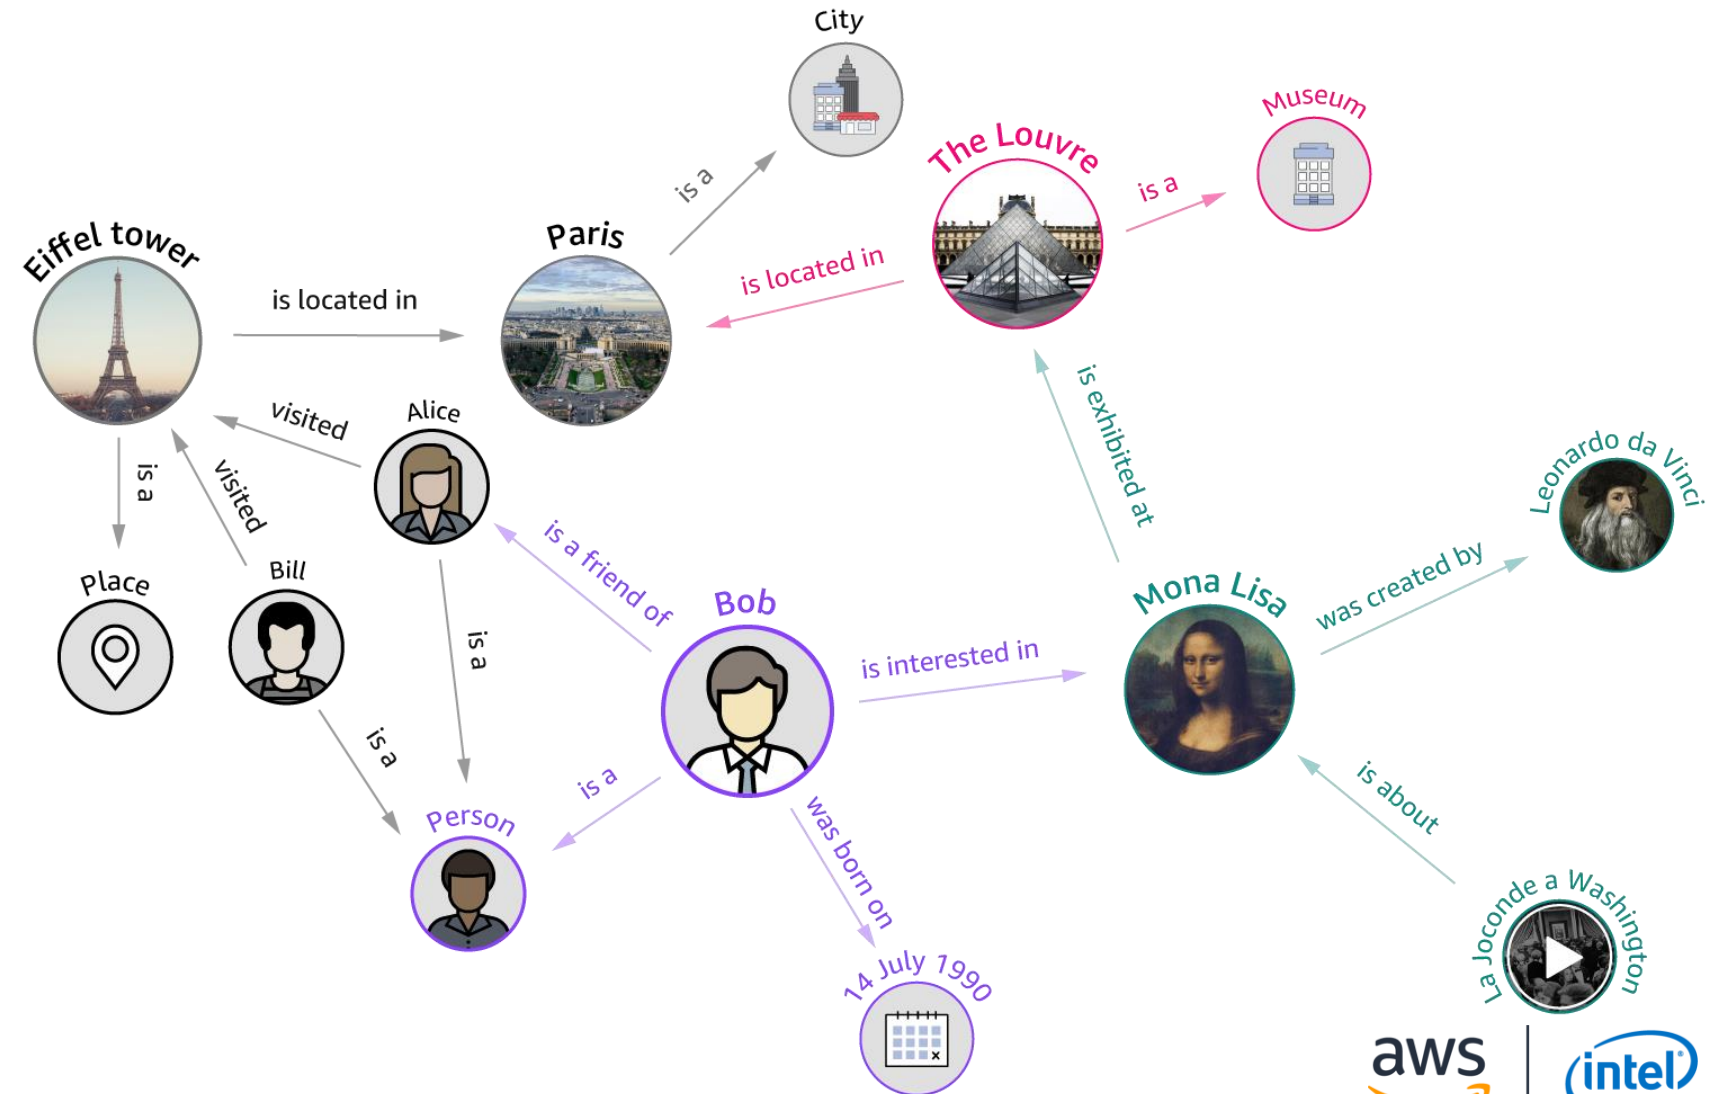


# Recommendations Based on Relationships



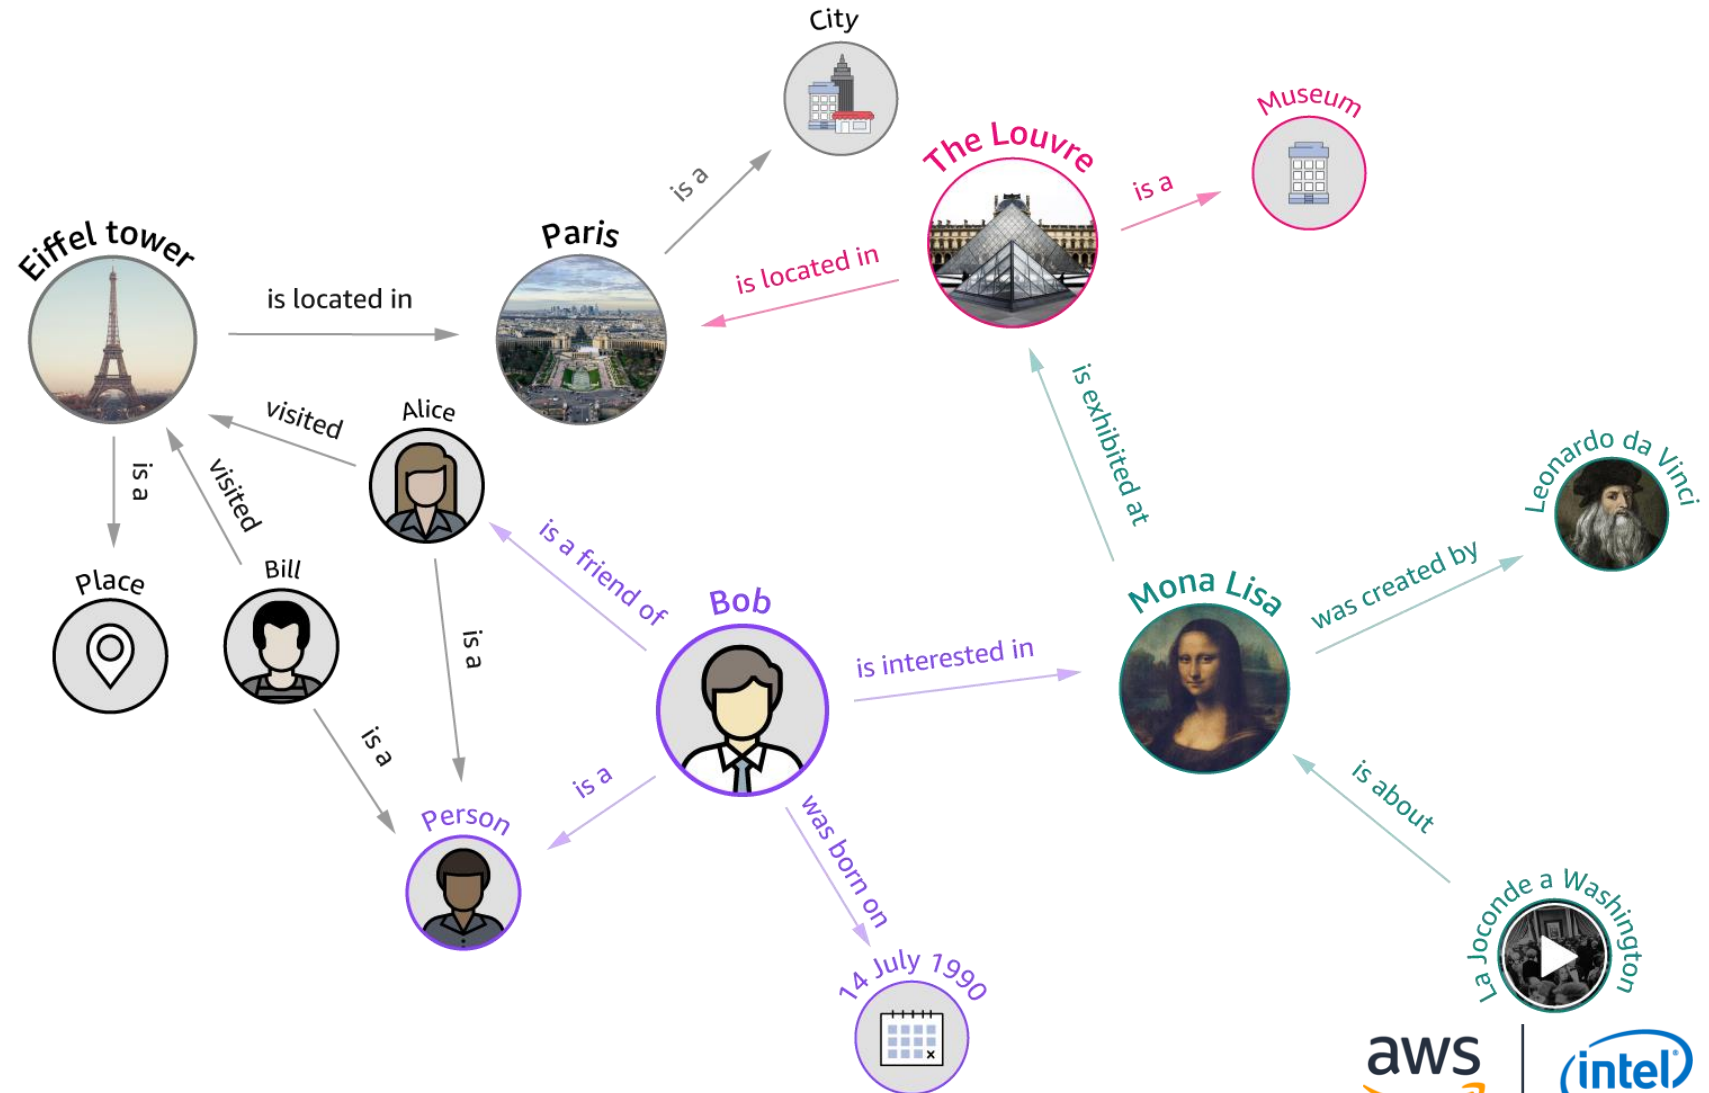


# Knowledge Graph Applications



# Knowledge Graph Applications

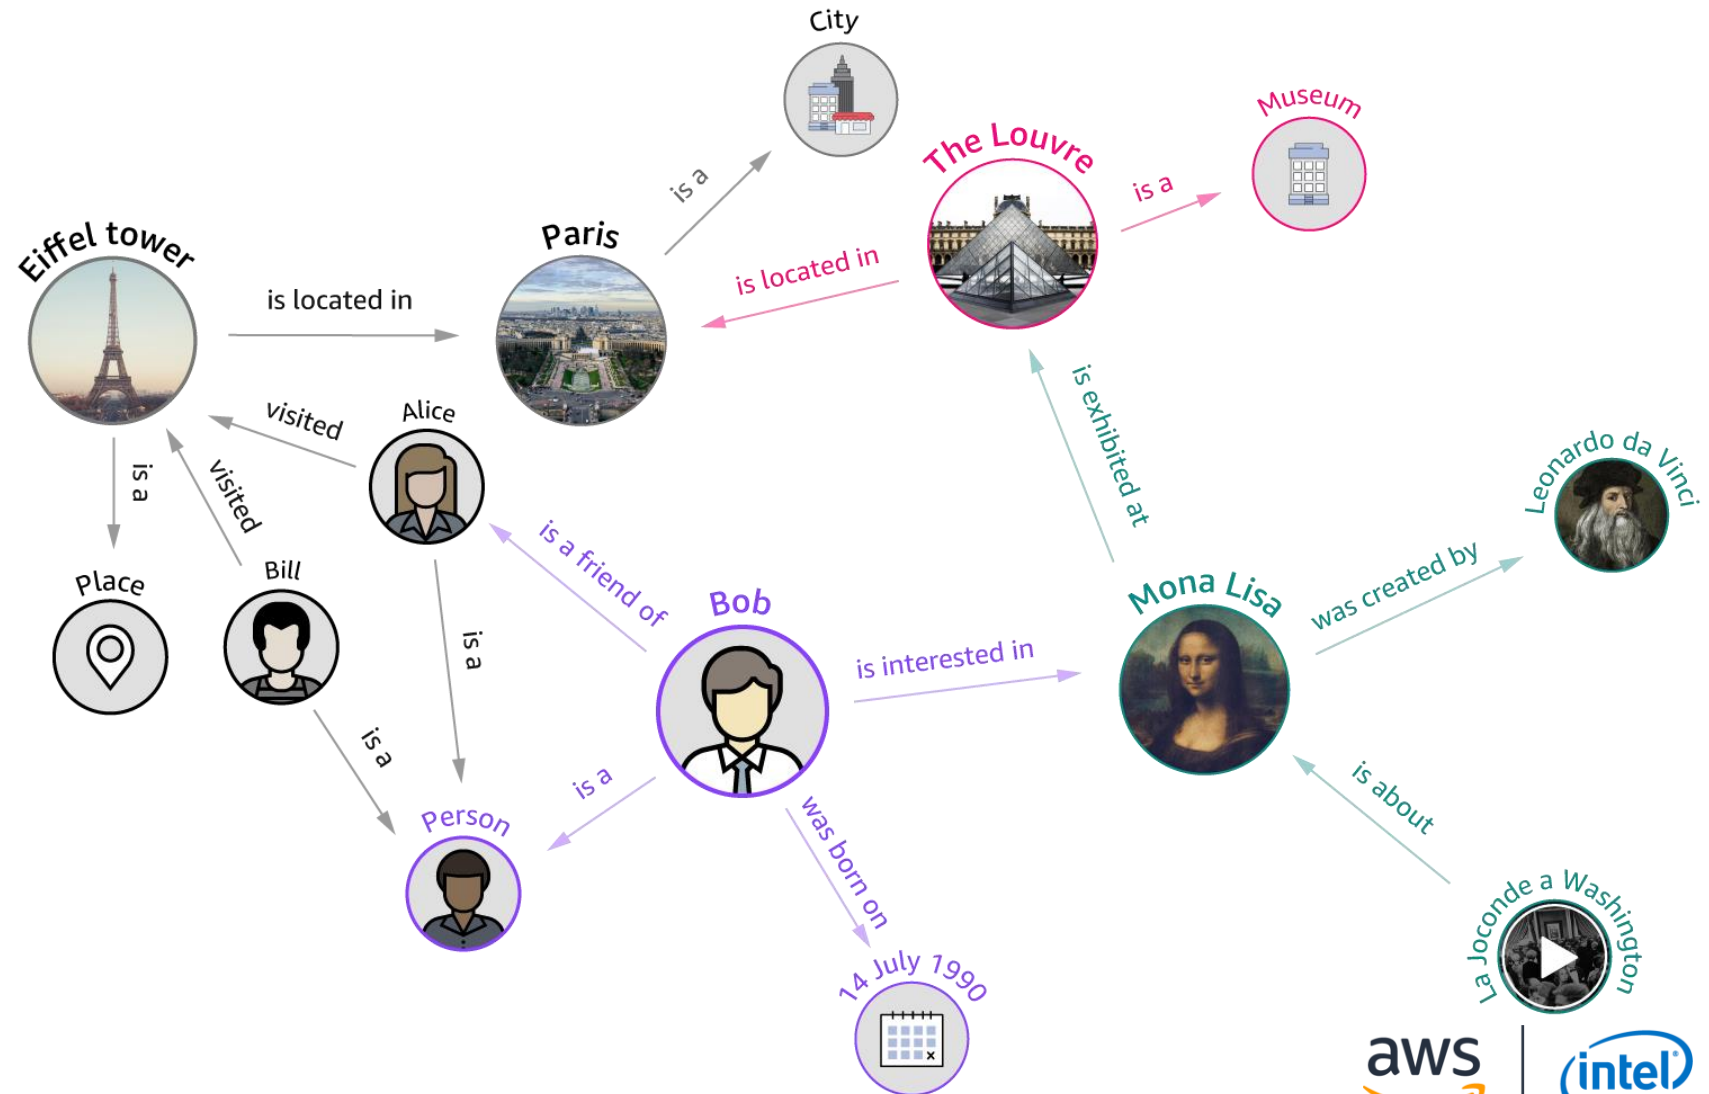
Who painted the Mona Lisa?



# Knowledge Graph Applications

Who painted the Mona Lisa?

What museums should Alice visit while in Paris?




# Knowledge Graph Applications

Who painted the Mona Lisa?

What museums should Alice visit while in Paris?

What artists have paintings in The Louvre?



# Navigate a Web of Global Tax Policies



THOMSON REUTERS


*"Our customers are increasingly required to navigate a complex web of global tax policies and regulations. We need an approach **to model the sophisticated corporate structures** of our largest clients and deliver an end-to-end tax solution. We use a microservices architecture approach for our platforms and are beginning to **leverage Amazon Neptune as a graph-based system** to quickly create links within the data."*

Tim Vanderham, Chief Technology Officer, Thomson Reuters Tax & Accounting


# Relational Database Challenges Building Apps with Highly Connected Data



Unnatural for  
querying graph

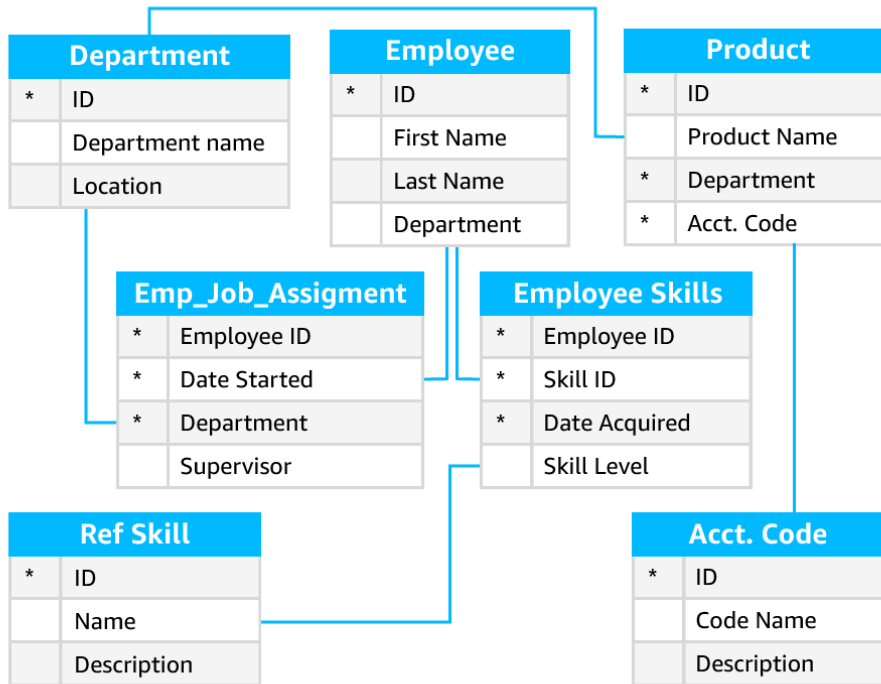


Inefficient  
graph processing

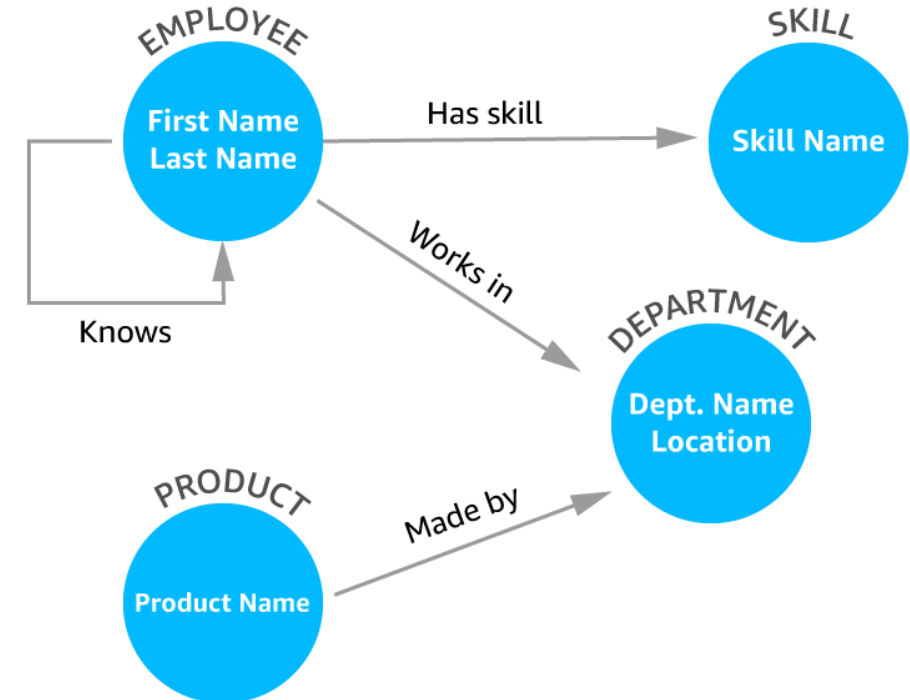


Rigid schema inflexible  
for changing data

# Different Approaches for Highly Connected Data




Purpose-built for a business process



Purpose-built to answer questions about relationships

# A Graph Database is Optimized for Efficient Storage and Retrieval of Highly Connected Data






# Leading Graph Models and Frameworks

## PROPERTY GRAPH

Open Source Apache TinkerPop™  
Gremlin Traversal Language



## RESOURCE DESCRIPTION FRAMEWORK (RDF)


W3C Standard  
SPARQL Query Language




# Challenges of Existing Graph Databases



Difficult to scale



Difficult to maintain  
high availability



Too expensive



Limited support for  
open standards

# Amazon Neptune - Fully Managed Graph Database

## OPEN




Supports Apache TinkerPop & W3C RDF graph models

## FAST



Query billions of relationships with millisecond latency

## RELIABLE



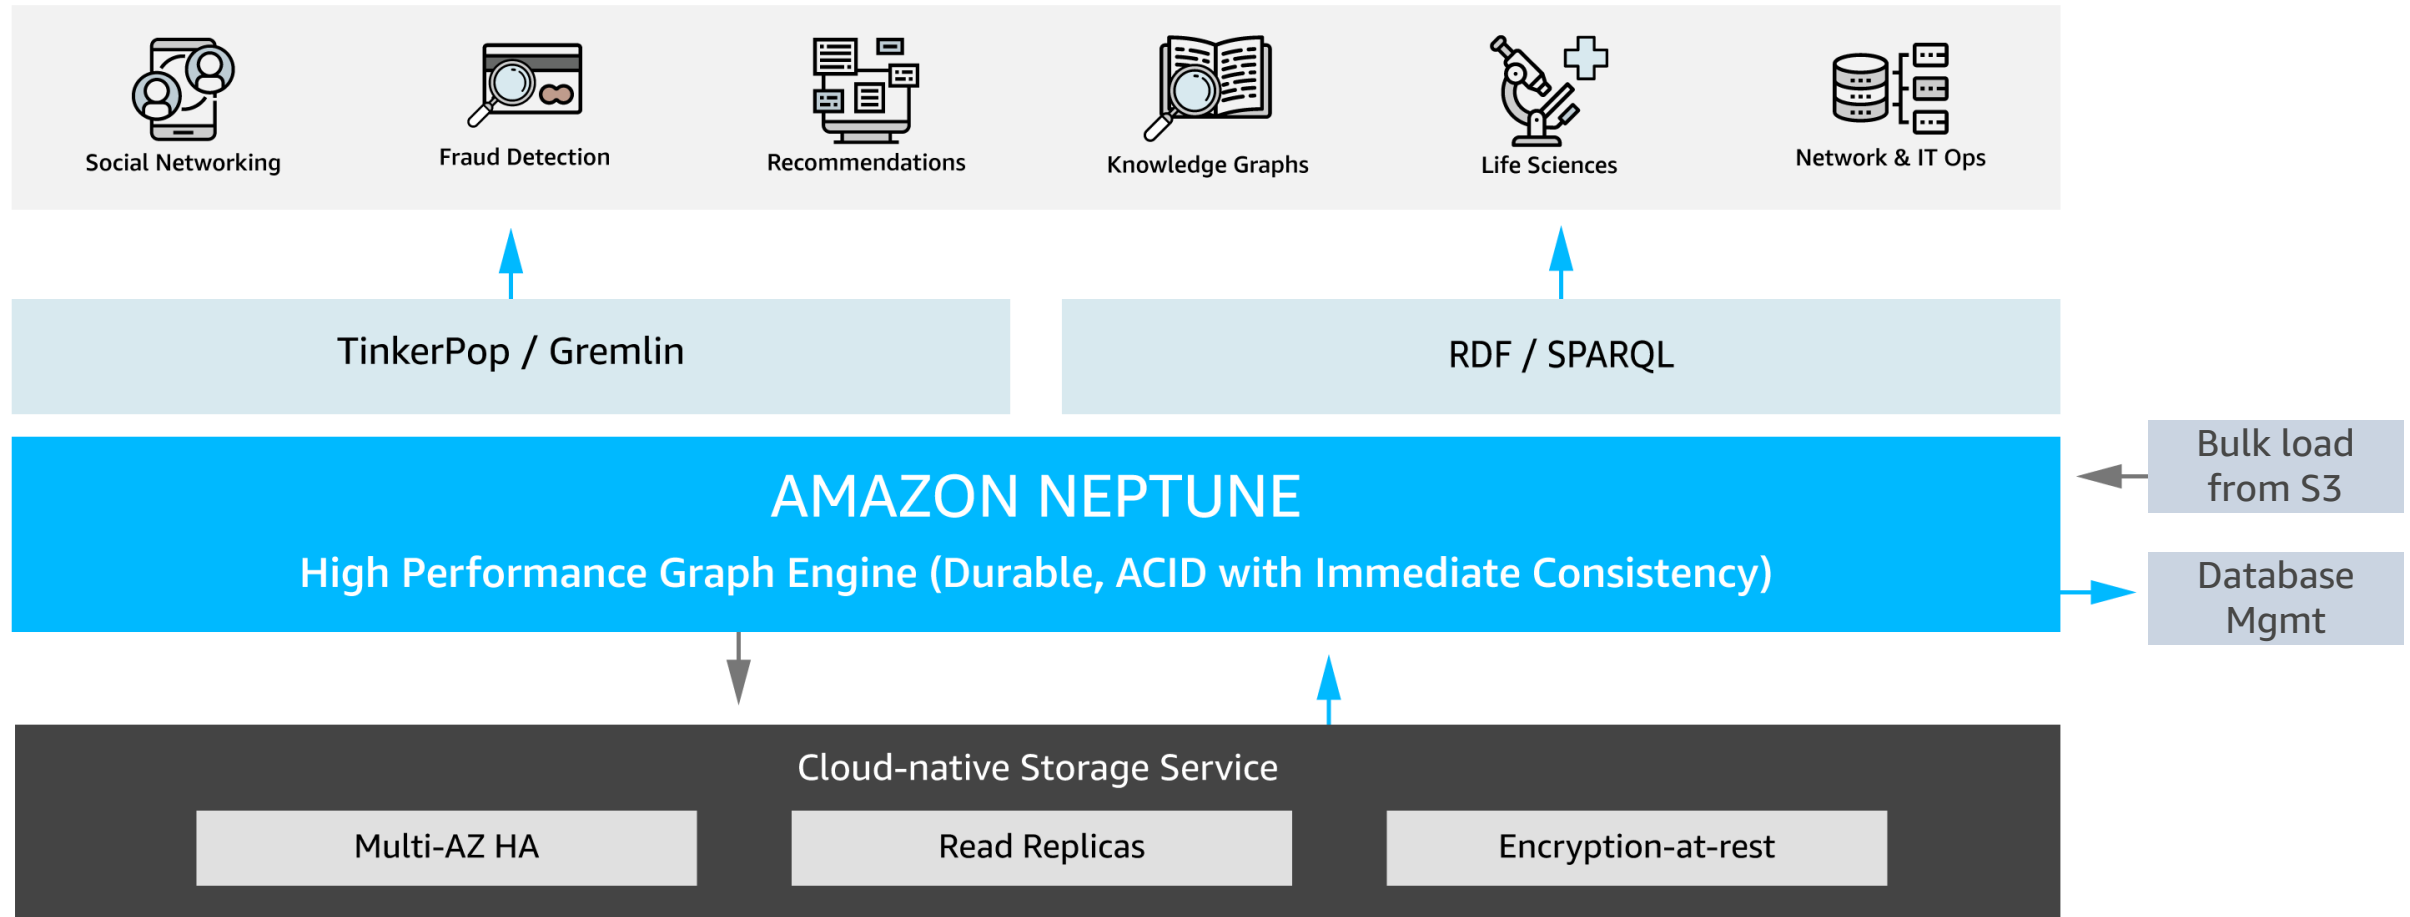
6 replicas of your data across 3 AZs with full backup and restore

## EASY



Build powerful queries easily with Gremlin and SPARQL


# Amazon Neptune High Level Architecture



# Types of Graphs and How to Query Them

## PROPERTY GRAPH

Open Source Apache TinkerPop™  
**Gremlin Traversal Language**




## RESOURCE DESCRIPTION FRAMEWORK (RDF)

W3C Standard  
**SPARQL Query Language**



# Property Graph

- A property graph is a set of vertices and edges with respective properties (i.e. key/value pairs)
- **Vertex** represents entities/domains
- **Edge** represents directional relationship between vertices.
  - Each edge has a label that denotes the type of relationship
- Each vertex & edge has a unique identifier
- Vertex and edges can have **properties**
- **Properties** express non-relational information about the vertices and edges



# Creating a TinkerPop Graph

//Connect to Neptune and receive a remote graph, g.

```
user1 = g.addVertex (id, 1, label, "User", "name", "Bill");
```


```
user2 = g.addVertex (id, 2, label, "User", "name", "Sarah");
```

...

```
user1.addEdge("FRIEND", user2, id, 21);
```




Gremlin (Apache TinkerPop 3.3)



# RDF Graphs

- **RDF Graphs** are described as a collection of triples: subject, predicate, and object.

subject      → `<http://www.socialnetwork.com/person#1>`  
predicate    →        `rdf:type contacts:User;`  
Object (literal) →        `contact:name: "Bill" .`



- **Internationalized Resource Identifiers (IRIs)** uniquely identify subjects.

IRI            → `<http://www.socialnetwork.com/person#1>`

- The Object can be an IRI or Literal.
  - A Literal in RDF is like a property and RDF supports the XML data types.
  - When the Object is an IRI, it forms an “Edge” in the graph.

subject      → `<http://www.socialnetwork.com/person#1>`  
predicate    →        `contacts:friend`  
Object (IRI) →        `<http://www.socialnetwork.com/person#2> .`





# “There’s No Trouble with Triples”: RDF Example

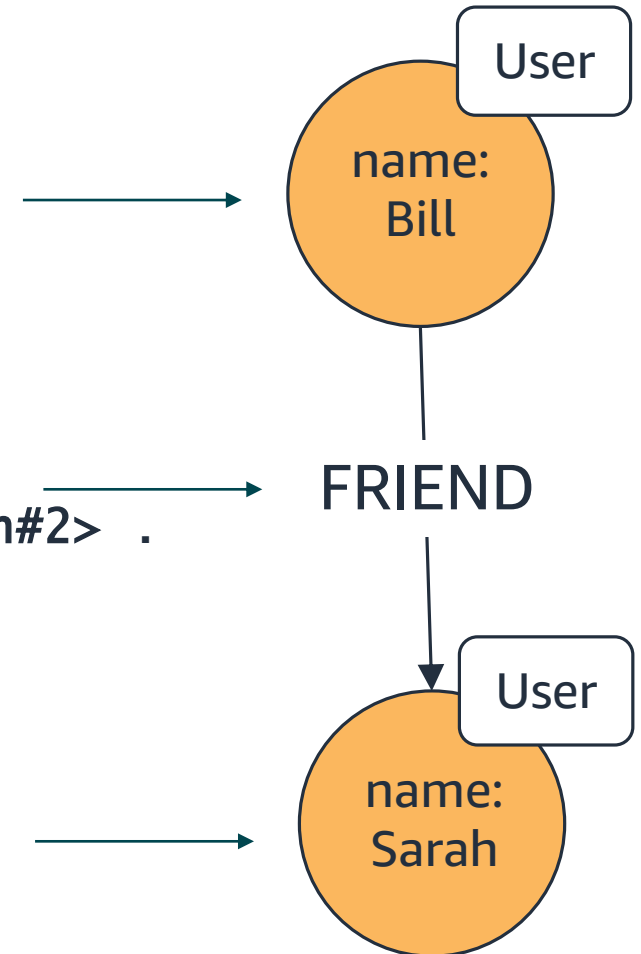
```
@prefix contacts: <http://www.socialnetwork.com/people#>.  
<http://www.socialnetwork.com/person#1>  
  rdf:type contacts:User;  
  contact:name: "Bill" .
```

```
<http://www.socialnetwork.com/person#1>  
  contacts:friend <http://www.socialnetwork.com/person#2> .
```

- <http://www.socialnetwork.com/person#2>
- rdf:type contacts:User;
- contact:name: "Sarah" .

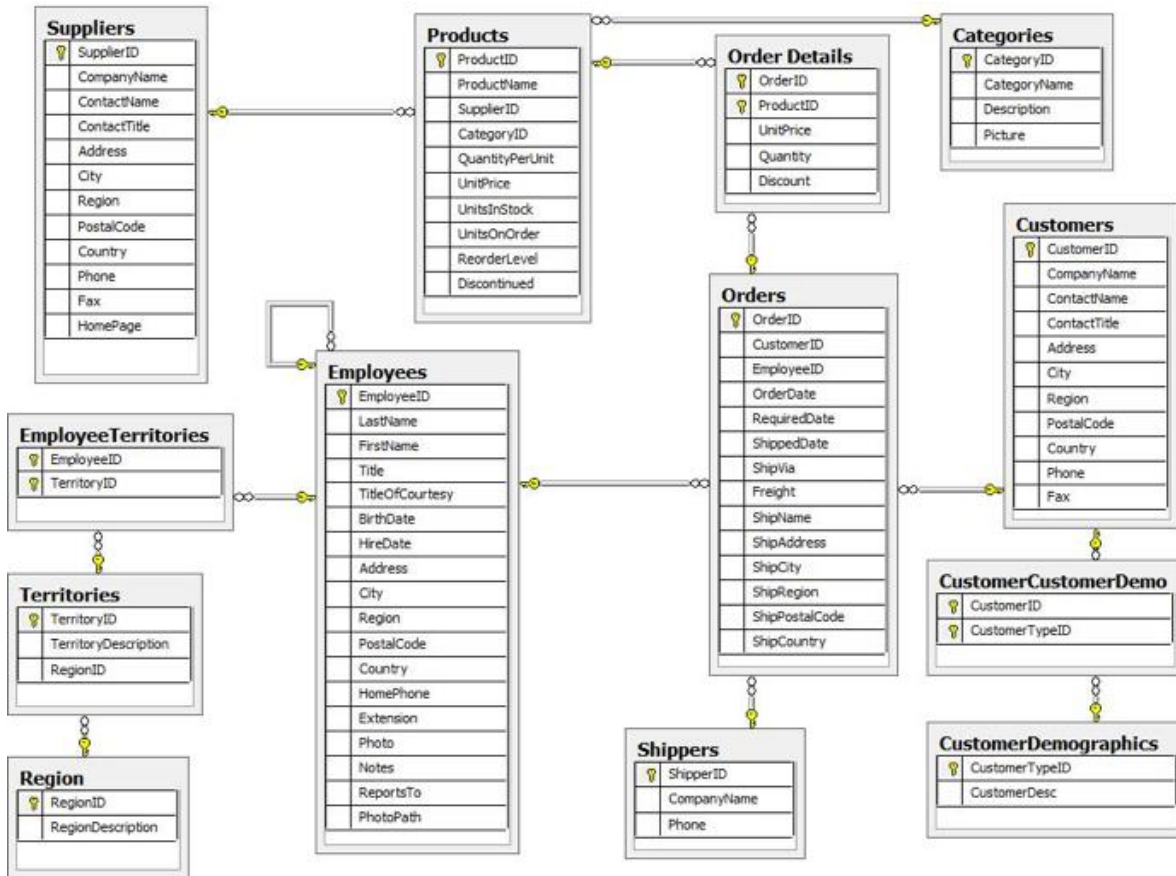


RDF  
(Turtle Serialization)

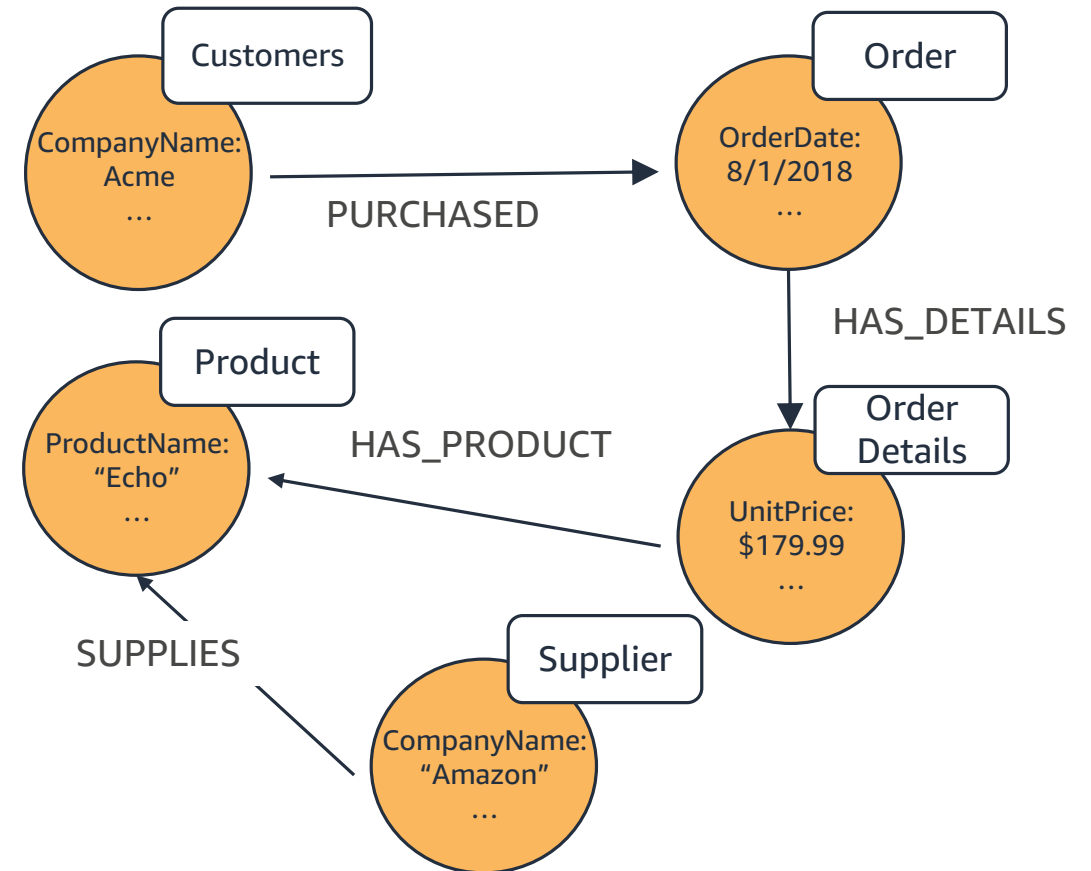


# Graph vs. Relational Database Modeling

## Relational model



## Graph model subset



\* Source : <http://www.playnexacro.com/index.html#show:article>



# SQL Relational Database Query

Find the name of companies that purchased the 'Echo'.

```
SELECT distinct c.CompanyName
```

```
FROM customers AS c
```

- **JOIN** orders AS o ON /\* Join the customer from  
the order \*/
- (c.CustomerID = o.CustomerID)
- **JOIN** order\_details AS od /\* Join the order details from the order \*/  
ON (o.OrderID = od.OrderID)
- **JOIN** products as p /\* Join the products from the order details \*/  
ON (od.ProductID = p.ProductID)
- **WHERE** p.ProductName = 'Echo'; /\* Find the product named 'Echo' \*/

# SPARQL Declarative Graph Query

Find the name of companies that purchased the 'Echo'.

```
PREFIX sales_db: <http://sales.widget.com/>
SELECT distinct ?comp_name WHERE {
  •      ?customer      <sales_db:HAS_ORDER> ?order ;           #customer graph pattern
  •      <sales_db:CompanyName> ?comp_name .                   #orders graph
  pattern
  •      ?order          <sales_db:HAS_DETAILS> ?order_d .     #order details graph
  pattern
  •      ?order_d        <sales_db:HAS_PRODUCT> ?product .      #products graph
  pattern
      ?product          <sales_db:ProductName> "Echo" .
}
```


# Gremlin Imperative Graph Traversal

Find the name of companies that purchased the 'Echo'.

```
/* All products named "Echo" */  
g.V().hasLabel('Product').has('name', 'Echo')  
  .in('HAS_PRODUCT') /* Traverse to order details */  
  .in('HAS_DETAILS') /* Traverse to order */  
  .in('HAS_ORDER') /* Traverse to Customer */  
  .values('CompanyName').dedup() /* Unique Company Name */
```

# TinkerPop Social Network Friend Recommendation Example

# Triadic Closure – Closing Triangles




# Recommending New Connections







# Immediate Friendships



# Means and Motive



# Recommendation



# Recommend New Connections

```
g = graph.traversal()

g.v().has('name', 'Terry').as('user').
  both('FRIEND').aggregate('friends').
  both('FRIEND').
    where(neq('user')).where(neq('friends')).
  groupCount().by('name').
  order(local).by(values, decr)
```

# Find Terry

```
g = graph.traversal()

g.v().has('name', 'Terry').as('user').
  both('FRIEND').aggregate('friends').
  both('FRIEND').
    where(neq('user')).where(neq('friends')).
  groupCount().by('name').
  order(local).by(values, decr)
```

# Find Terry's Friends

```
g = graph.traversal()


g.V().has('name', 'Terry').as('user').
  both('FRIEND').aggregate('friends').
  both('FRIEND').
    where(neq('user')).where(neq('friends')).
  groupCount().by('name').
  order(local).by(values, decr)
```

# And The Friends of Those Friends

```
g = graph.traversal()
```

```
g.V().has('name', 'Terry').as('user').  
  both('FRIEND').aggregate('friends').  
both('FRIEND').
```

```
  where(neq('user')).where(neq('friends'))  
  groupCount().by('name').  
  order(local).by(values, decr)
```




# ...Who Aren't Terry and Aren't Friends with Terry

```
g = graph.traversal()
```

```
g.V().has('name', 'Terry').as('user').  
  both('FRIEND').aggregate('friends').  
  both('FRIEND').
```

```
  where(neq('user')).where(neq('friends'))
```

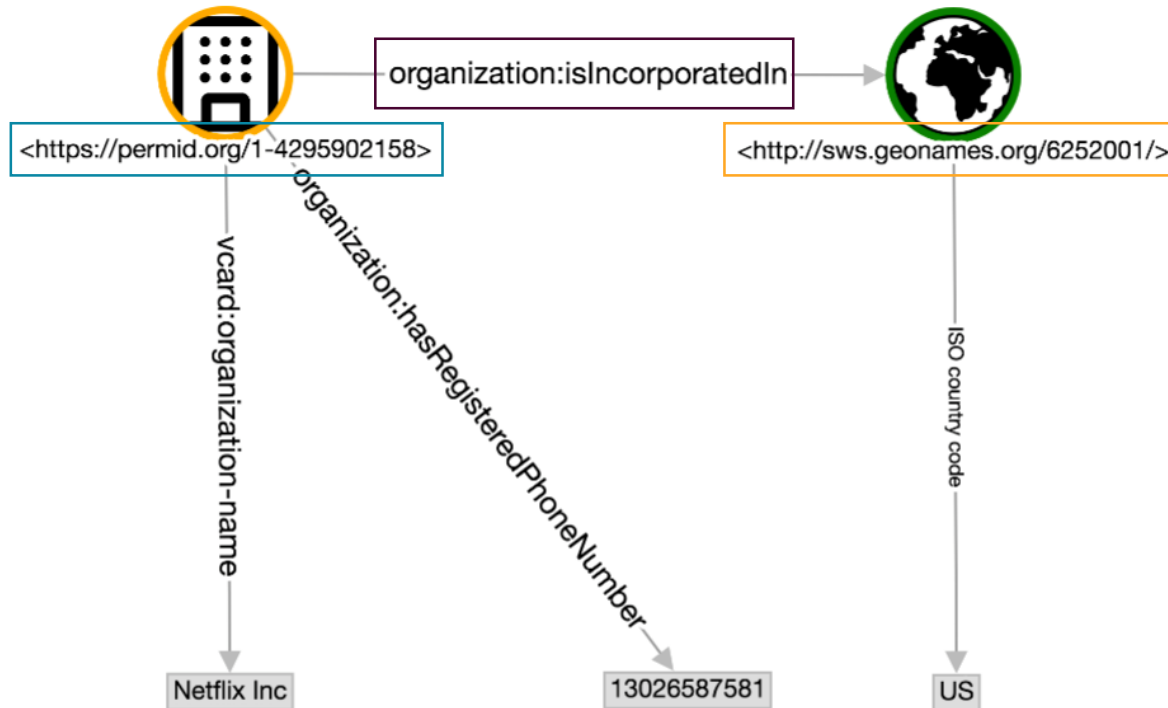
```
  groupCount().by('name').  
  order(local).by(values, decr)
```





# RDF Knowledge Graph Example

# URIs as Globally Unique Identifiers



## URIs to identify nodes and edge labels

`<https://permid.org/1-4295902158>`

=> identifies the company "Netflix Inc"

`organization:isIncorporatedIn1`

=> identifies the relationship "is incorporated in"

`<http://sws.geonames.org/6252001/>`

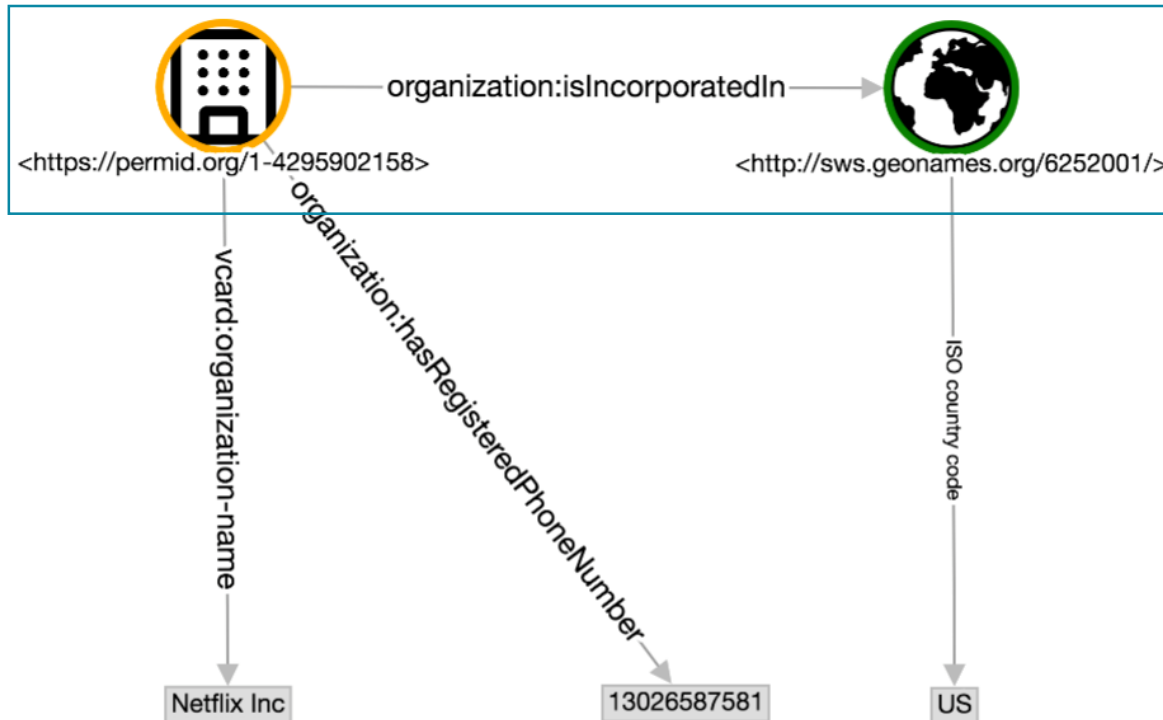
=> identifies country "USA"

<sup>1</sup> This is a shortcut for

`<http://permid.org/ontology/organization/isIncorporatedIn>`.

RDF uses XML prefix notation, where the prefix `organization` is a shortcut for `<http://permid.org/ontology/organization/>`.

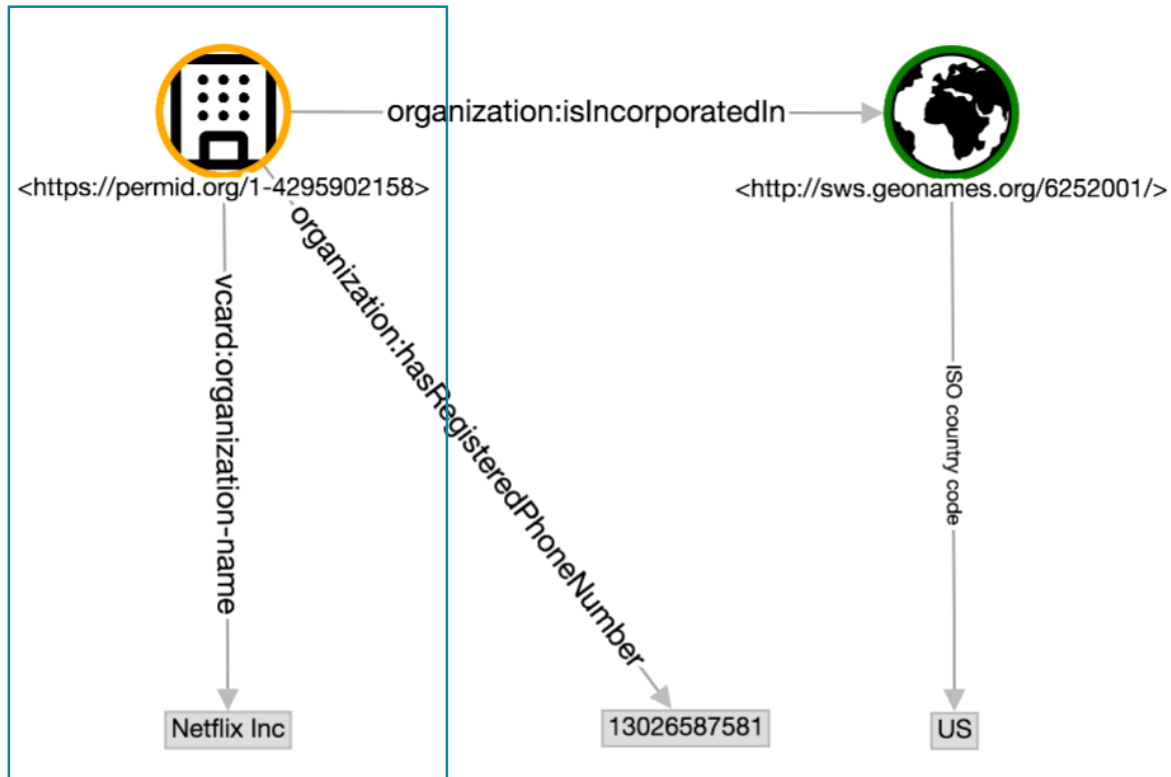
# RDF Graph: Collection of Triples (1)



- # Every edge in the RDF graph is represented as
- # a (subject, predicate, object) triple

- |   |             |
|---|-------------|
| • <code>&lt;https://permid.org/1-4295902158&gt;</code>  | ← subject   |
| • <code>organization:isIncorporatedIn</code>            | ← predicate |
| • <code>&lt;http://sws.geonames.org/6252001/&gt;</code> | ← object    |

# RDF Graph: Collection of Triples (2)



- # Every edge in the RDF graph is represented as
- # a (subject, predicate, object) triple

- `<https://permid.org/1-4295902158>` ← subject
- `organization:isIncorporatedIn` ← predicate
- `<http://sws.geonames.org/6252001/>` ← object (URI)

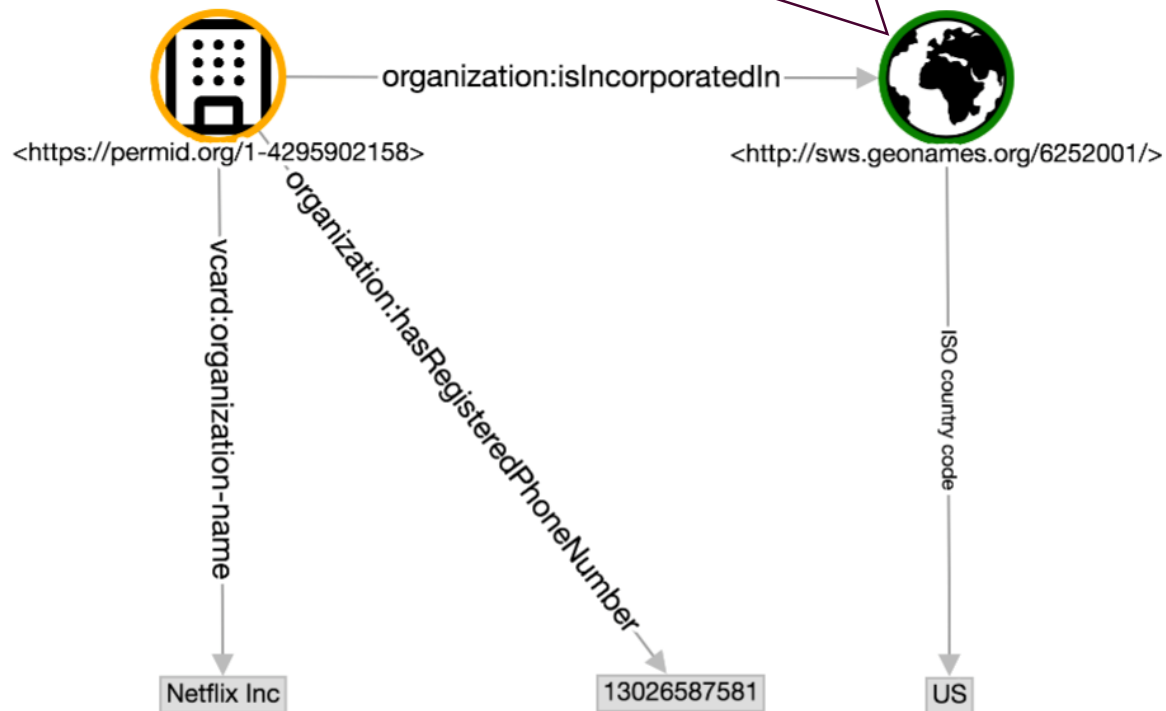
- `<https://permid.org/1-4295902158>` ← subject
- `vcard:organization-name` ← predicate
- `"Netflix Inc"` . ← object (literal)

Literals are "sinks" in the graph. They do not have any outgoing edges.

RDF supports strings and other XML datatypes (bool, integer, dates, floats, doubles, ...)

# RDF Graph: Collection of Triples (3)

Same URI used as both subject and object, depending on whether we represent outgoing vs. incoming RDF edges.



- # Every edge in the RDF graph is represented as
- # a (subject, predicate, object) triple

- `<https://permid.org/1-4295902158>`
- `organization:isIncorporatedIn`
- `<http://sws.geonames.org/6252001/>` .

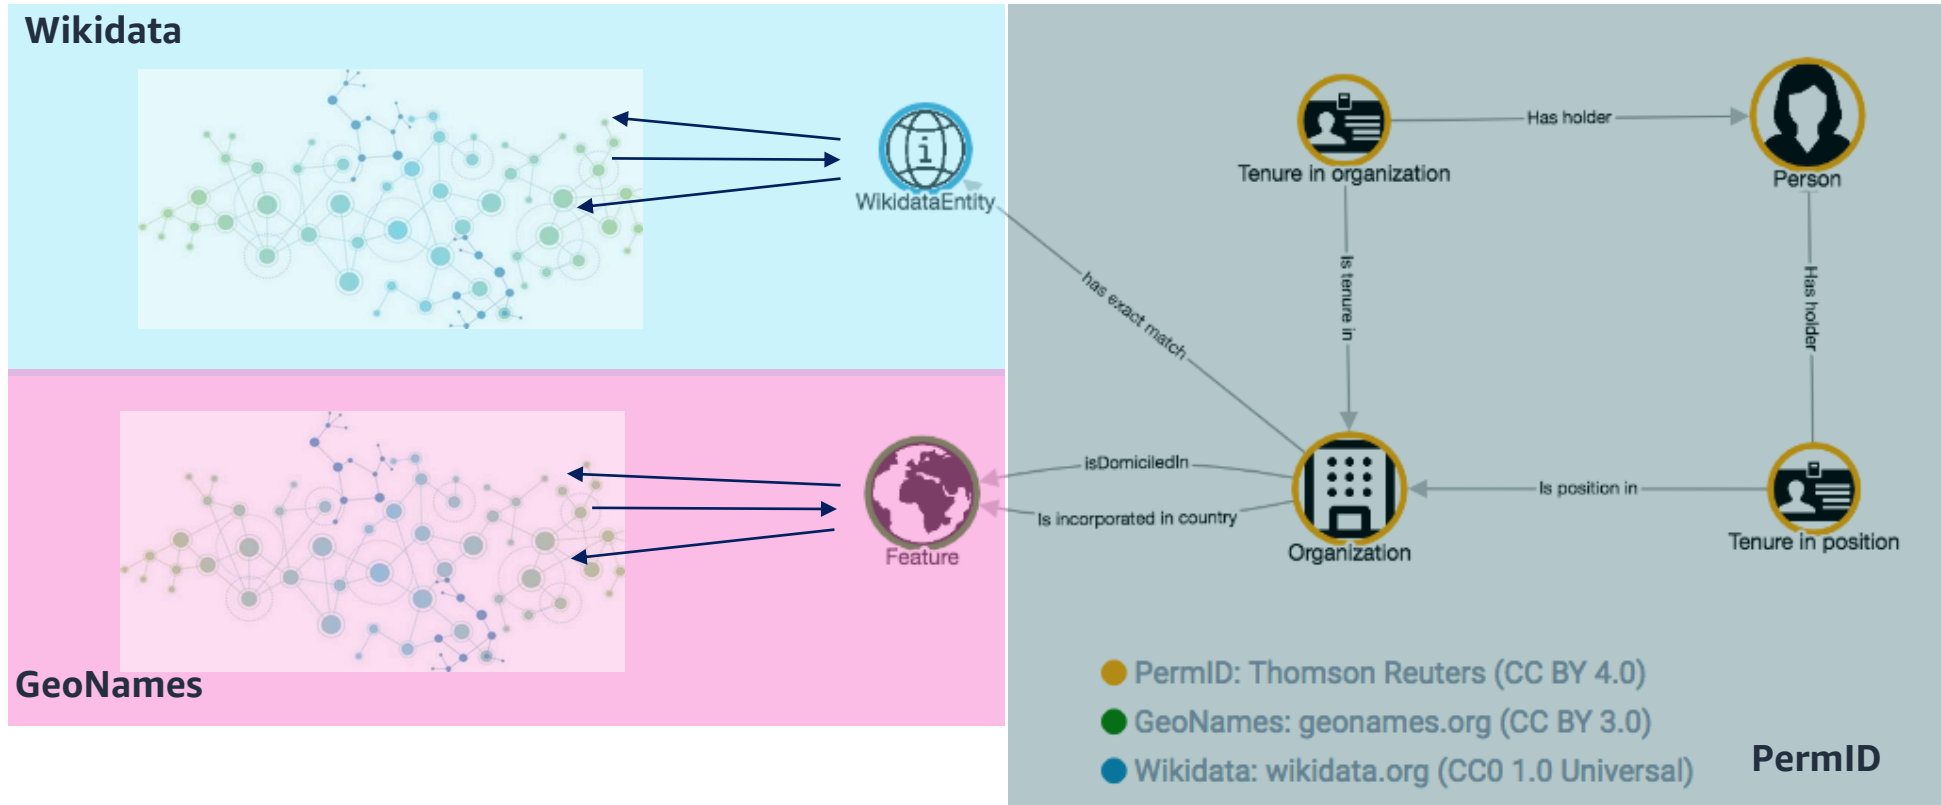
- `<https://permid.org/1-4295902158>`
- `vcard:organization-name`
- `"Netflix Inc"` .

- `<https://permid.org/1-4295902158>`
- `organization:hasRegisteredPhoneNumber`
- `"13026587581"` .

- `<http://sws.geonames.org/6252001/>`
- `iso:countryCode`
- `"US"` .

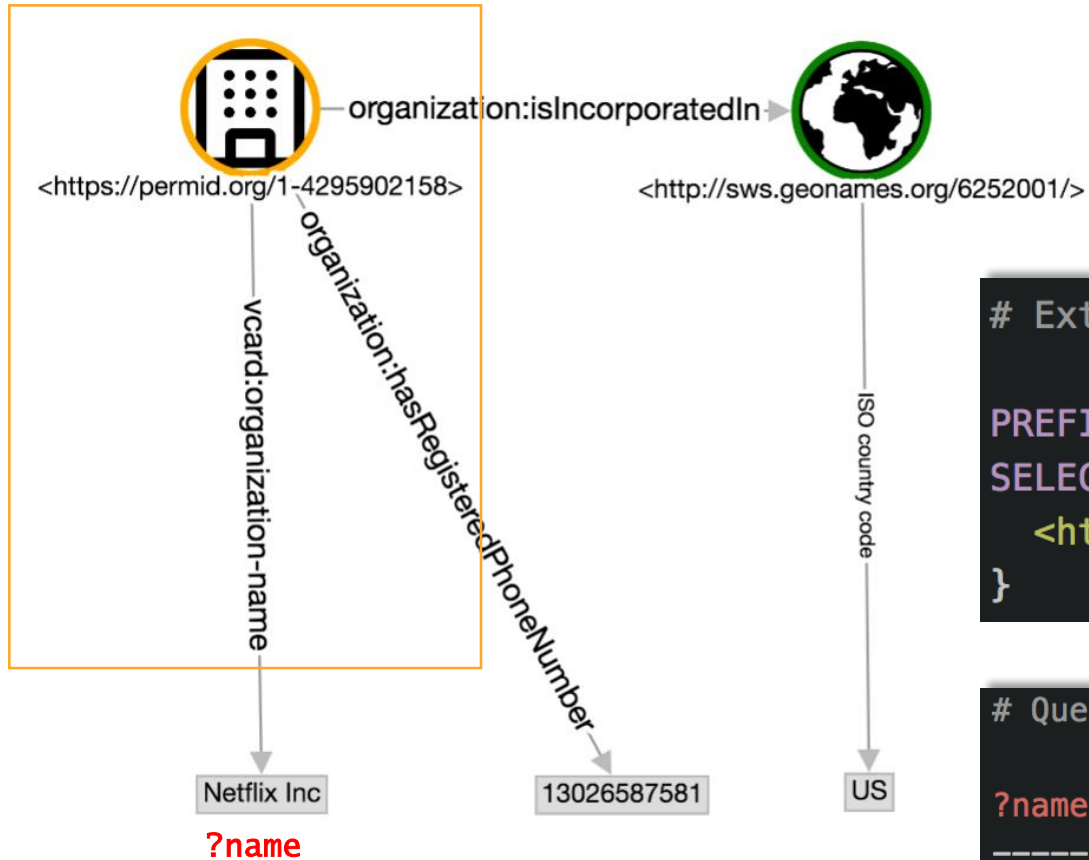
# The Benefit of URIs: Linked Data

Linking across datasets by referencing globally unique URIs



Example: PermID (re)uses `<http://sws.geonames.org/6252001/>` as a global Identifier for the USA, which is an identifier rooted in GeoNames.

# Querying RDF Using SPARQL (1)



Variables are prefixed with "?".

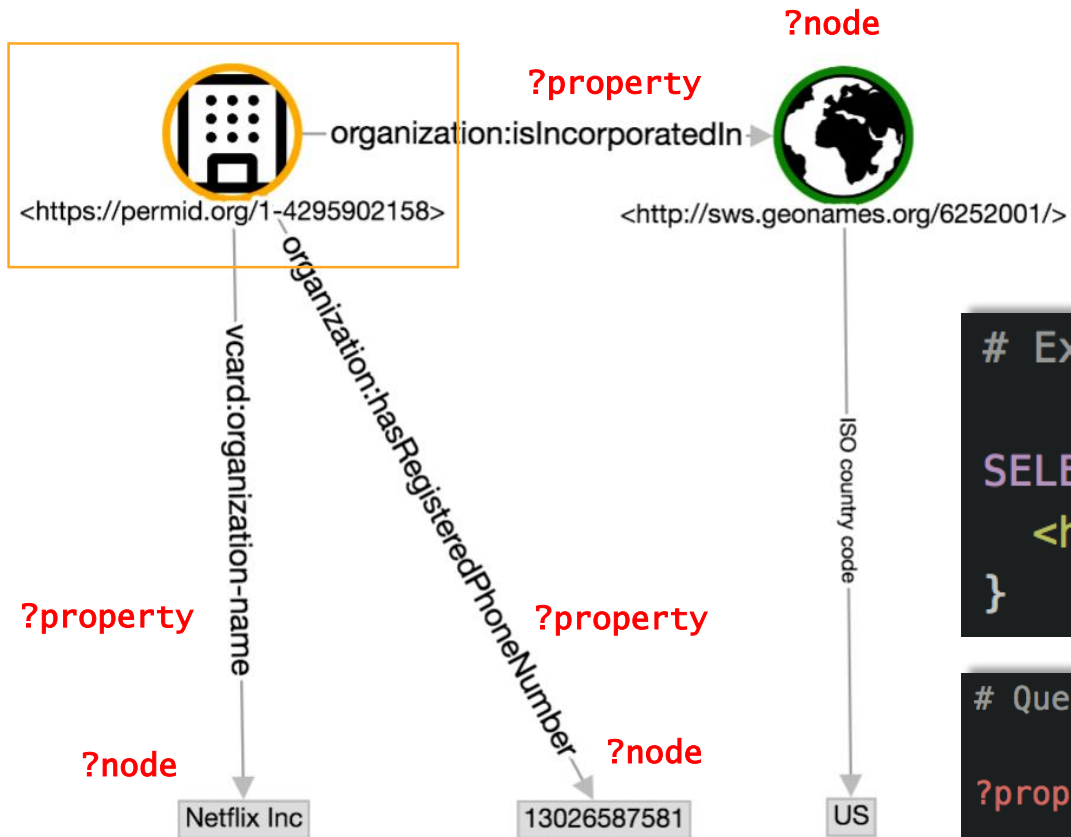
```
# Extract the name of company Netflix (using its URI)
```

```
PREFIX vcard: <http://www.w3.org/2006/vcard/ns#>
SELECT ?name WHERE {
  <https://permid.org/1-4295902158> vcard:organization-name ?name
}
```

```
# Query result:
```

```
?name
-----
"Netflix Inc"
```

# Querying RDF Using SPARQL (2)



# Extract outgoing predicate-to-object pairs for Netflix

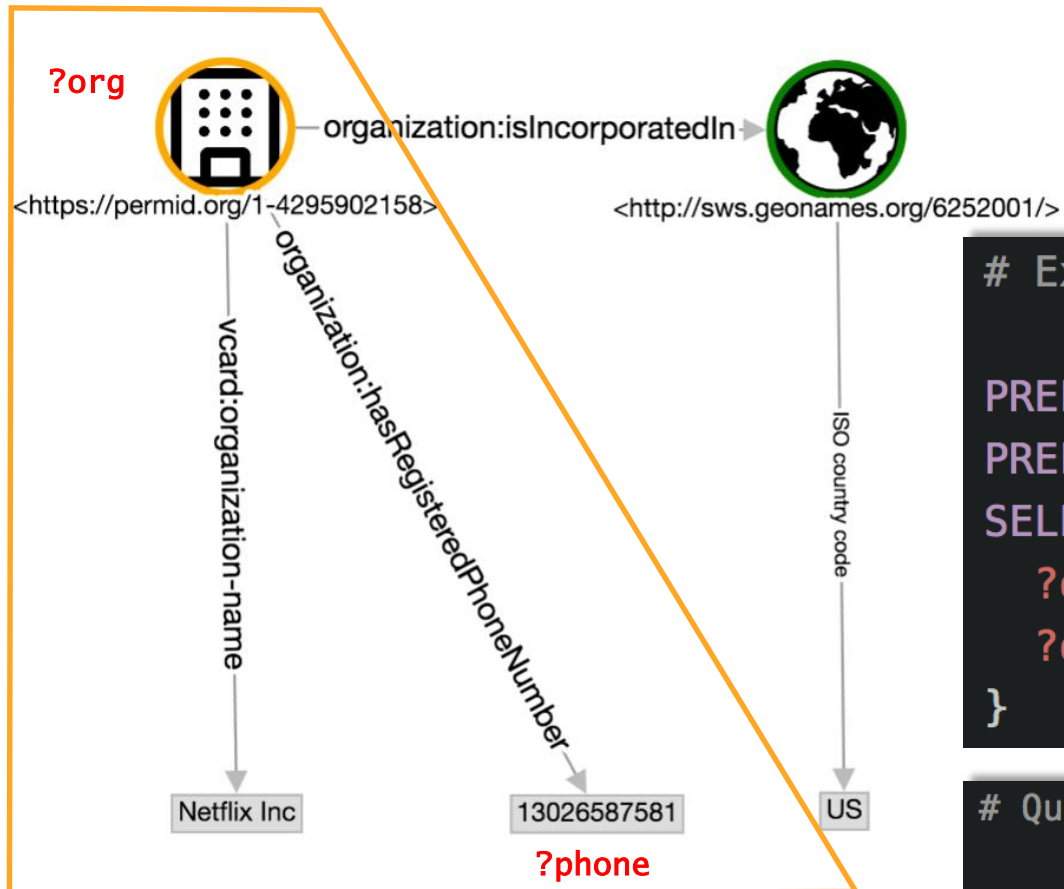
```
SELECT ?property ?node WHERE {  
  <https://permid.org/1-4295902158> ?property ?node  
}
```

# Query result:

?property	?node
<code>vcard:organization-name</code>	<code>"Netflix Inc"</code>
<code>organization:hasRegisteredPhoneNumber</code>	<code>"13026587581"</code>
<code>organization:isIncorporatedIn</code>	<code>&lt;http://sws/geo...&gt;</code>



# Querying RDF Using SPARQL (3)



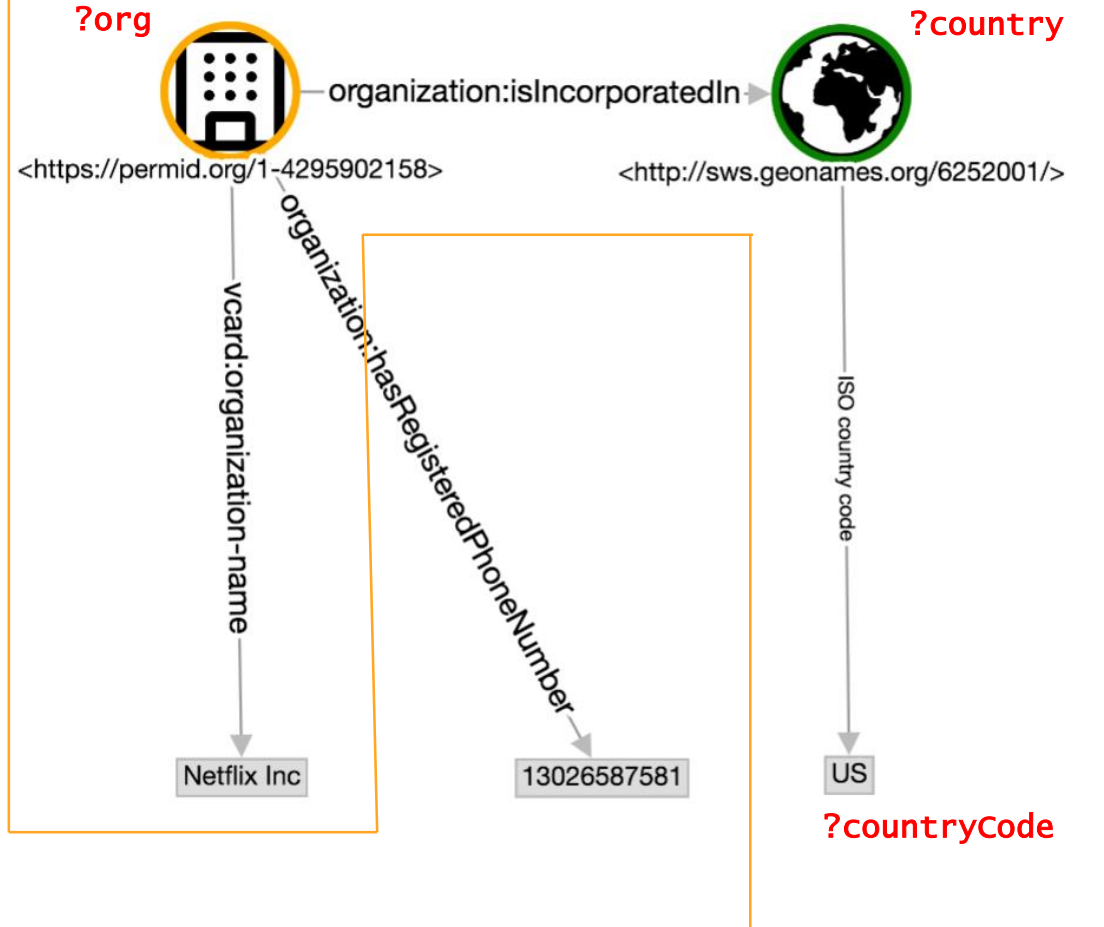
```
# Extract URI and phone number of company "Netflix Inc"
```

```
PREFIX vcard: <http://www.w3.org/2006/vcard/ns#>
PREFIX org: <http://permid.org/ontology/organization/>
SELECT ?org ?phone WHERE {
    ?org vcard:organization-name "Netflix Inc" .
    ?org org:hasRegisteredPhoneNumber ?phone
}
```

```
# Query result:
```

<code>?org</code>	<code>?phone</code>
<code>&lt;https://permid.org/1-4295902158&gt;</code>	<code>"13026587581"</code>

# Querying RDF Using SPARQL (4)



```
# Extract the organization ID and country code of  
# company "Netflix Inc"
```

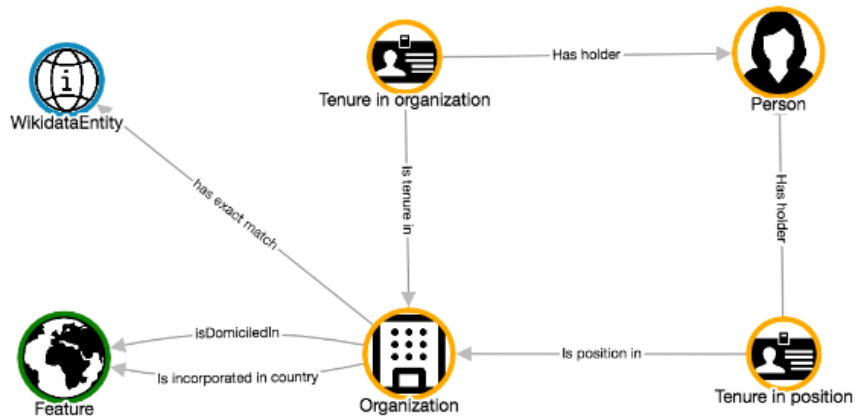
```
PREFIX vcard: <http://www.w3.org/2006/vcard/ns#>  
PREFIX org: <http://permid.org/ontology/organization/>  
PREFIX geo: <http://www.geonames.org/ontology#>
```

```
SELECT ?org ?countryCode WHERE {  
  ?org vcard:organization-name "Netflix Inc" .  
  ?org org:isIncorporatedIn ?country .  
  ?country geo:countryCode ?countryCode  
}
```

```
# Query result:
```

?org	?countryCode
<code>&lt;https://permid.org/1-4295902158&gt;</code>	<code>"US"</code>

# The Power of Linked Data



- PermID: Thomson Reuters (CC BY 4.0)
- GeoNames: geonames.org (CC BY 3.0)
- Wikidata: wikidata.org (CC0 1.0 Universal)

Data from PermID

Data from Wikidata

Data from GeoNames

```
SELECT DISTINCT ?companyName ?companyId
  (SAMPLE(?stockExchangeWD) AS ?sampleStockExchangePerWD)
  (MAX(?employeeNumberWD) AS ?maxEmployeeNumberPerWD)
  ?locationEnglishName ?locationChineseName
```

```
WHERE {
```

```
# extract company name, Wikidata URI, and location from PermID
```

```
?companyId vcard:organization-name ?companyName .
```

```
?companyId skos:exactMatch ?orgInWikidata .
```

```
?companyId fibo:isDomiciledIn ?location .
```

```
# filter by companies with more than one hundred thousand employees
```

```
# and extract stock exchange information from Wikidata
```

```
?orgInWikidata <http://www.wikidata.org/prop/P1128> ?employeeStmt .
```

```
?employeeStmt <http://www.wikidata.org/prop/statement/P1128> ?employeeNumberWD .
```

```
FILTER (?employeeNumberWD > 10000)
```

```
?orgInWikidata <http://www.wikidata.org/prop/direct/P414> ?stockExchange .
```

```
?stockExchange rdfs:label ?stockExchangeWD .
```

```
# extract the company location's English and Chinese name
```

```
?location geo:alternateName ?locationEnglishName .
```

```
FILTER((LANG(?locationEnglishName)) = "en")
```

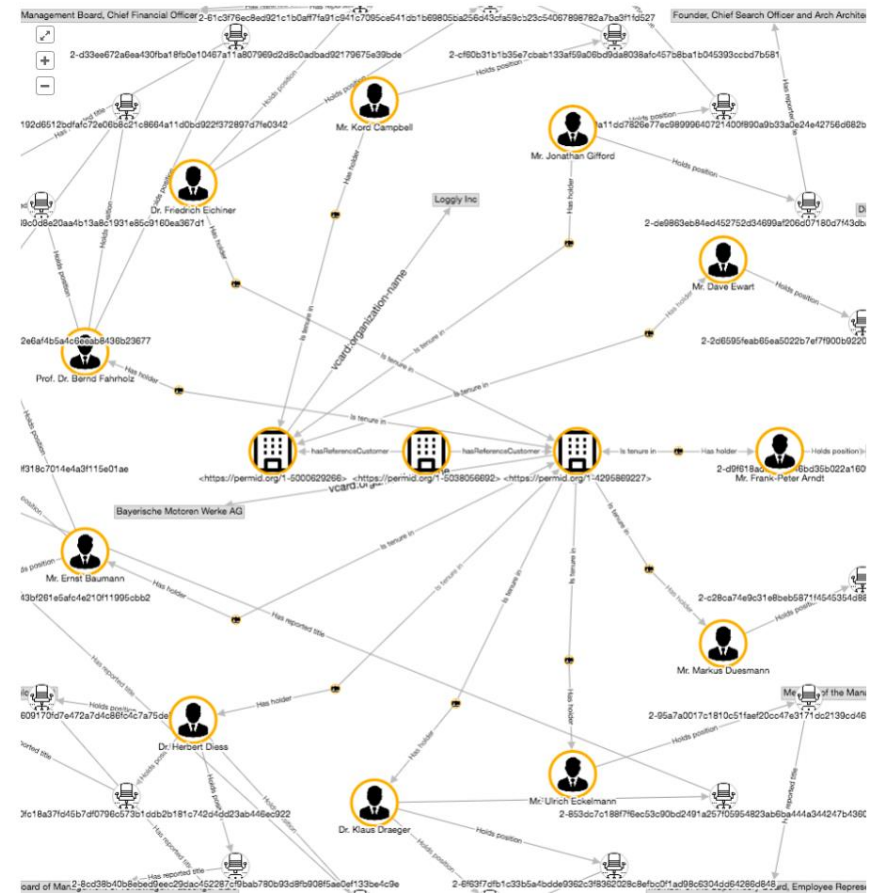
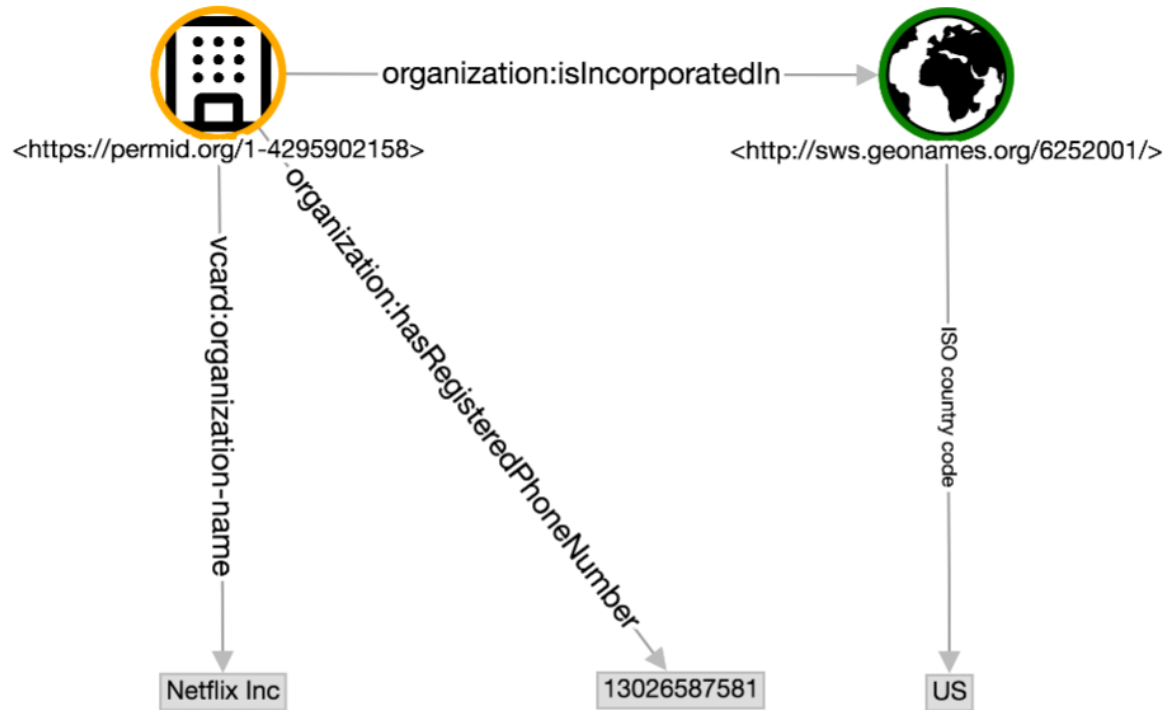
```
?location geo:alternateName ?locationChineseName .
```

```
FILTER((LANG(?locationChineseName)) = "zh")
```

```
} GROUP BY ?companyName ?companyId ?location
```

```
  ?locationEnglishName ?locationChineseName
```

# RDF Graphs

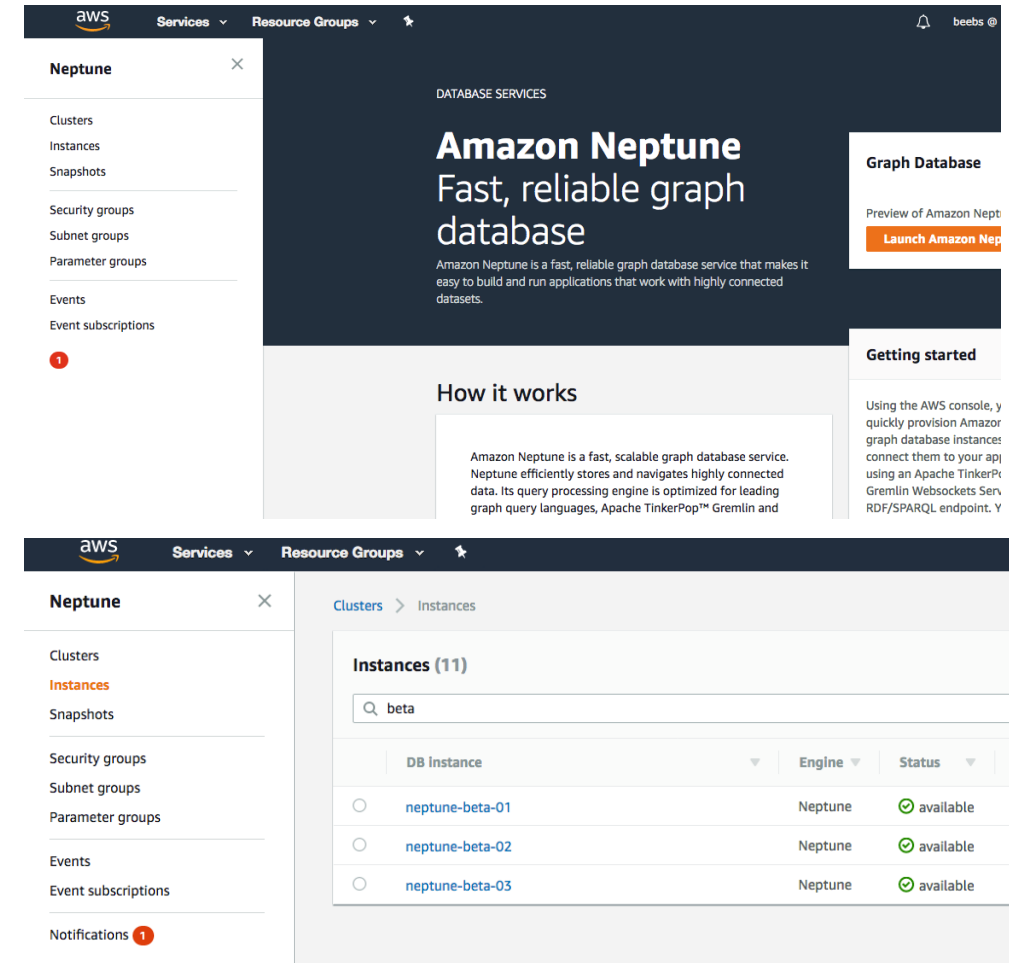


# Fully Managed Service

## BENEFITS

- ✓ Easily configurable via the Console
- ✓ Multi-AZ High Availability, ACID
- ✓ Support for up to 15 read replicas
- ✓ Supports Encryption at rest
- ✓ Supports Encryption in transit (TLS)
- ✓ Backup and Restore, Point-in-time Recovery

<https://aws.amazon.com/neptune/>



# Learn from AWS experts. Advance your skills and knowledge. Build your future in the AWS Cloud.



## Digital Training

Free, self-paced online courses built by AWS experts



## Classroom Training

Classes taught by accredited AWS instructors



## AWS Certification

Exams to validate expertise with an industry-recognized credential

Ready to begin building your cloud skills?  
Get started at: <https://www.aws.training/>

# With deep expertise on AWS, APN Partners can help your organization at any stage of your Cloud Adoption Journey.



## **AWS Managed Service Providers**

APN Consulting Partners who are skilled at cloud infrastructure and application migration, and offer proactive management of their customer's environment.



## **AWS Competency Partners**

APN Partners who have demonstrated technical proficiency and proven customer success in specialized solution areas.



## **AWS Marketplace**

A digital catalog with thousands of software listings from independent software vendors that make it easy to find, test, buy, and deploy software that runs on AWS.



## **AWS Service Delivery Partners**

APN Partners with a track record of delivering specific AWS services to customers.

**Ready to get started with an APN Partner?**  
**Find a partner:** <https://aws.amazon.com/partners/find/>  
**Learn more at the AWS Partner Network Booth**



# Thank You for Attending AWS Innovate

We hope you found it interesting! A kind reminder to **complete the survey.**

Let us know what you thought of today's event and how we can improve the event experience for you in the future.



[aws-apac-marketing@amazon.com](mailto:aws-apac-marketing@amazon.com)



[twitter.com/AWSCloud](https://twitter.com/AWSCloud)



[facebook.com/AmazonWebServices](https://facebook.com/AmazonWebServices)



[youtube.com/user/AmazonWebServices](https://youtube.com/user/AmazonWebServices)



[slideshare.net/AmazonWebServices](https://slideshare.net/AmazonWebServices)



[twitch.tv/aws](https://twitch.tv/aws)